

# Deploying Containerized Workloads Using Google Cloud Kubernetes Engine

---

INTRODUCING GOOGLE KUBERNETES ENGINE (GKE)



**Janani Ravi**

CO-FOUNDER, LOONYCORN

[www.loonycorn.com](http://www.loonycorn.com)



# Overview

**Containers for lightweight compute**

**Ideal for hybrid, multi-cloud**

**Kubernetes container orchestration technology**

**Industry standard with Google origins**

**GKE for Kubernetes on GCP**



# Prerequisites and Course Outline

---

# Software and Skills




**Basic understanding of cloud computing**

**Basic understanding of how virtual machines work**

**Basic understanding of how containers work**

# Prerequisites: Basic Cloud Computing



## Choosing and Implementing Google Cloud Compute Engine Solutions

- Basics of using the Google Cloud Platform

## Docker and Containers: The Big Picture

- Introduction to containers





# Course Outline

## Introducing GKE

- VMs vs. containers
- Terminology: Pods, kubelets, node pools clusters

## Creating and administering clusters

- Clusters using the web console and command line
- Autoscaling, autohealing and autorepairing nodes

## Deploying containerized workloads

- Deploying containers and exposing services
- Sharing state with persistent volumes
- User requests with ingress objects
- Binary authorization for attestations

## Monitoring clusters with Stackdriver

- Stackdriver Kubernetes monitoring and Prometheus



# Scenarios: SpikeySales.com



## Hypothetical Online Retailer

- Flash sales of trending products
- Spikes in user traffic

## SpikeySales on the GCP

- Cloud computing fits perfectly
- Pay-as-you-go
- No idle capacity during off-sale periods

## SpikeySales and Cloud Storage Buckets

- Move data to the cloud
- Elastic, pay-as-you-go, global access

# Introducing Containers

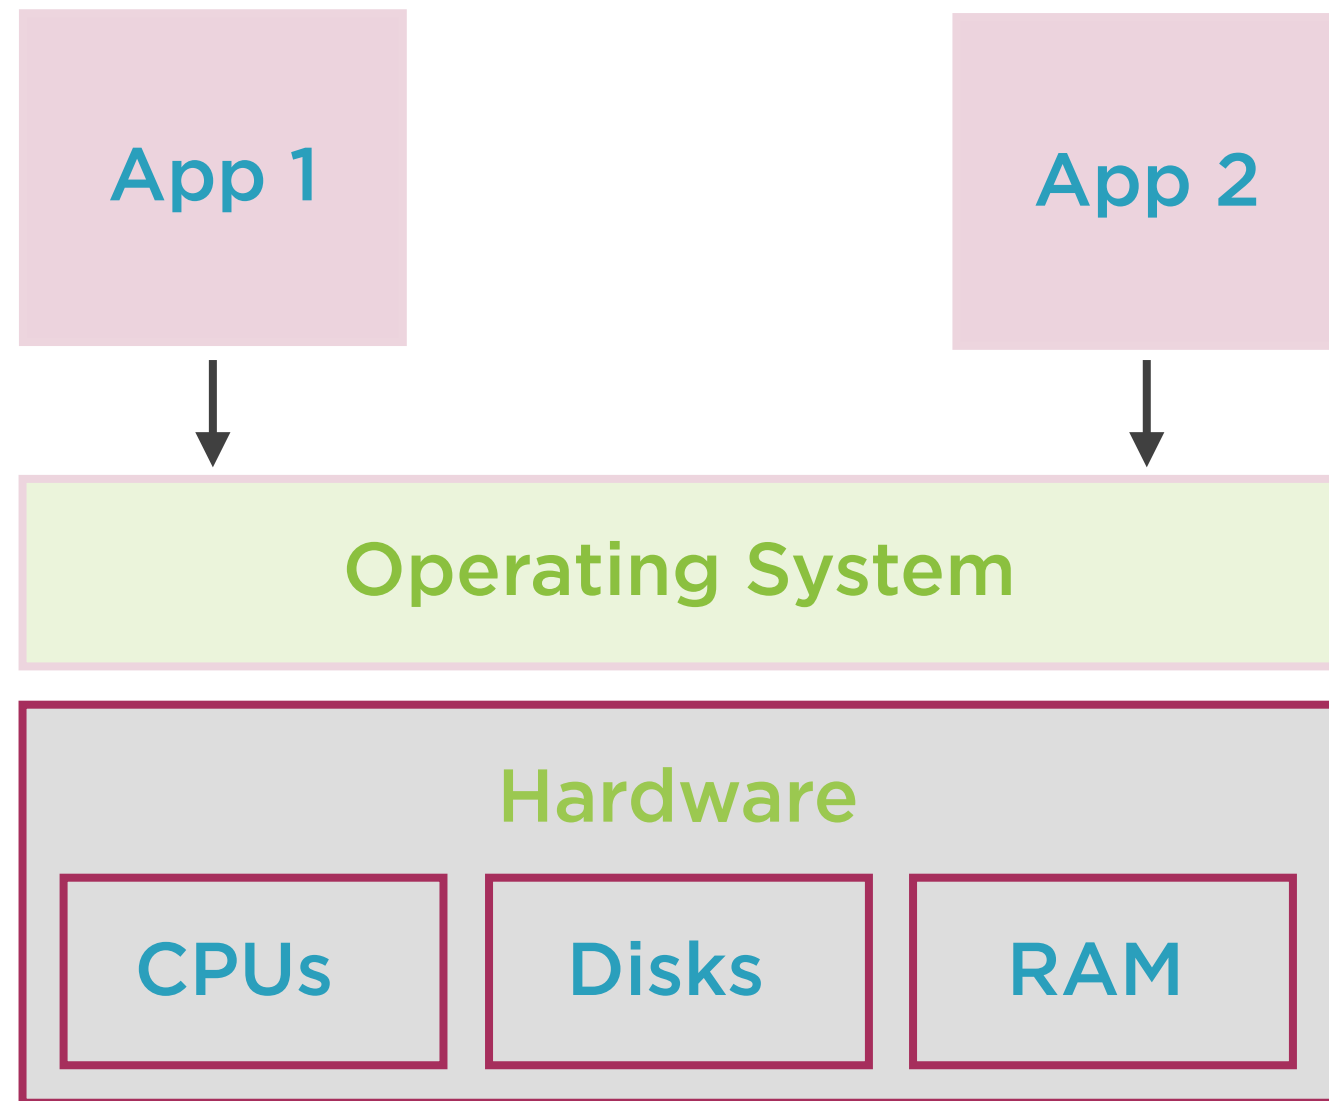
---






# Bare Metal and Virtual Machines

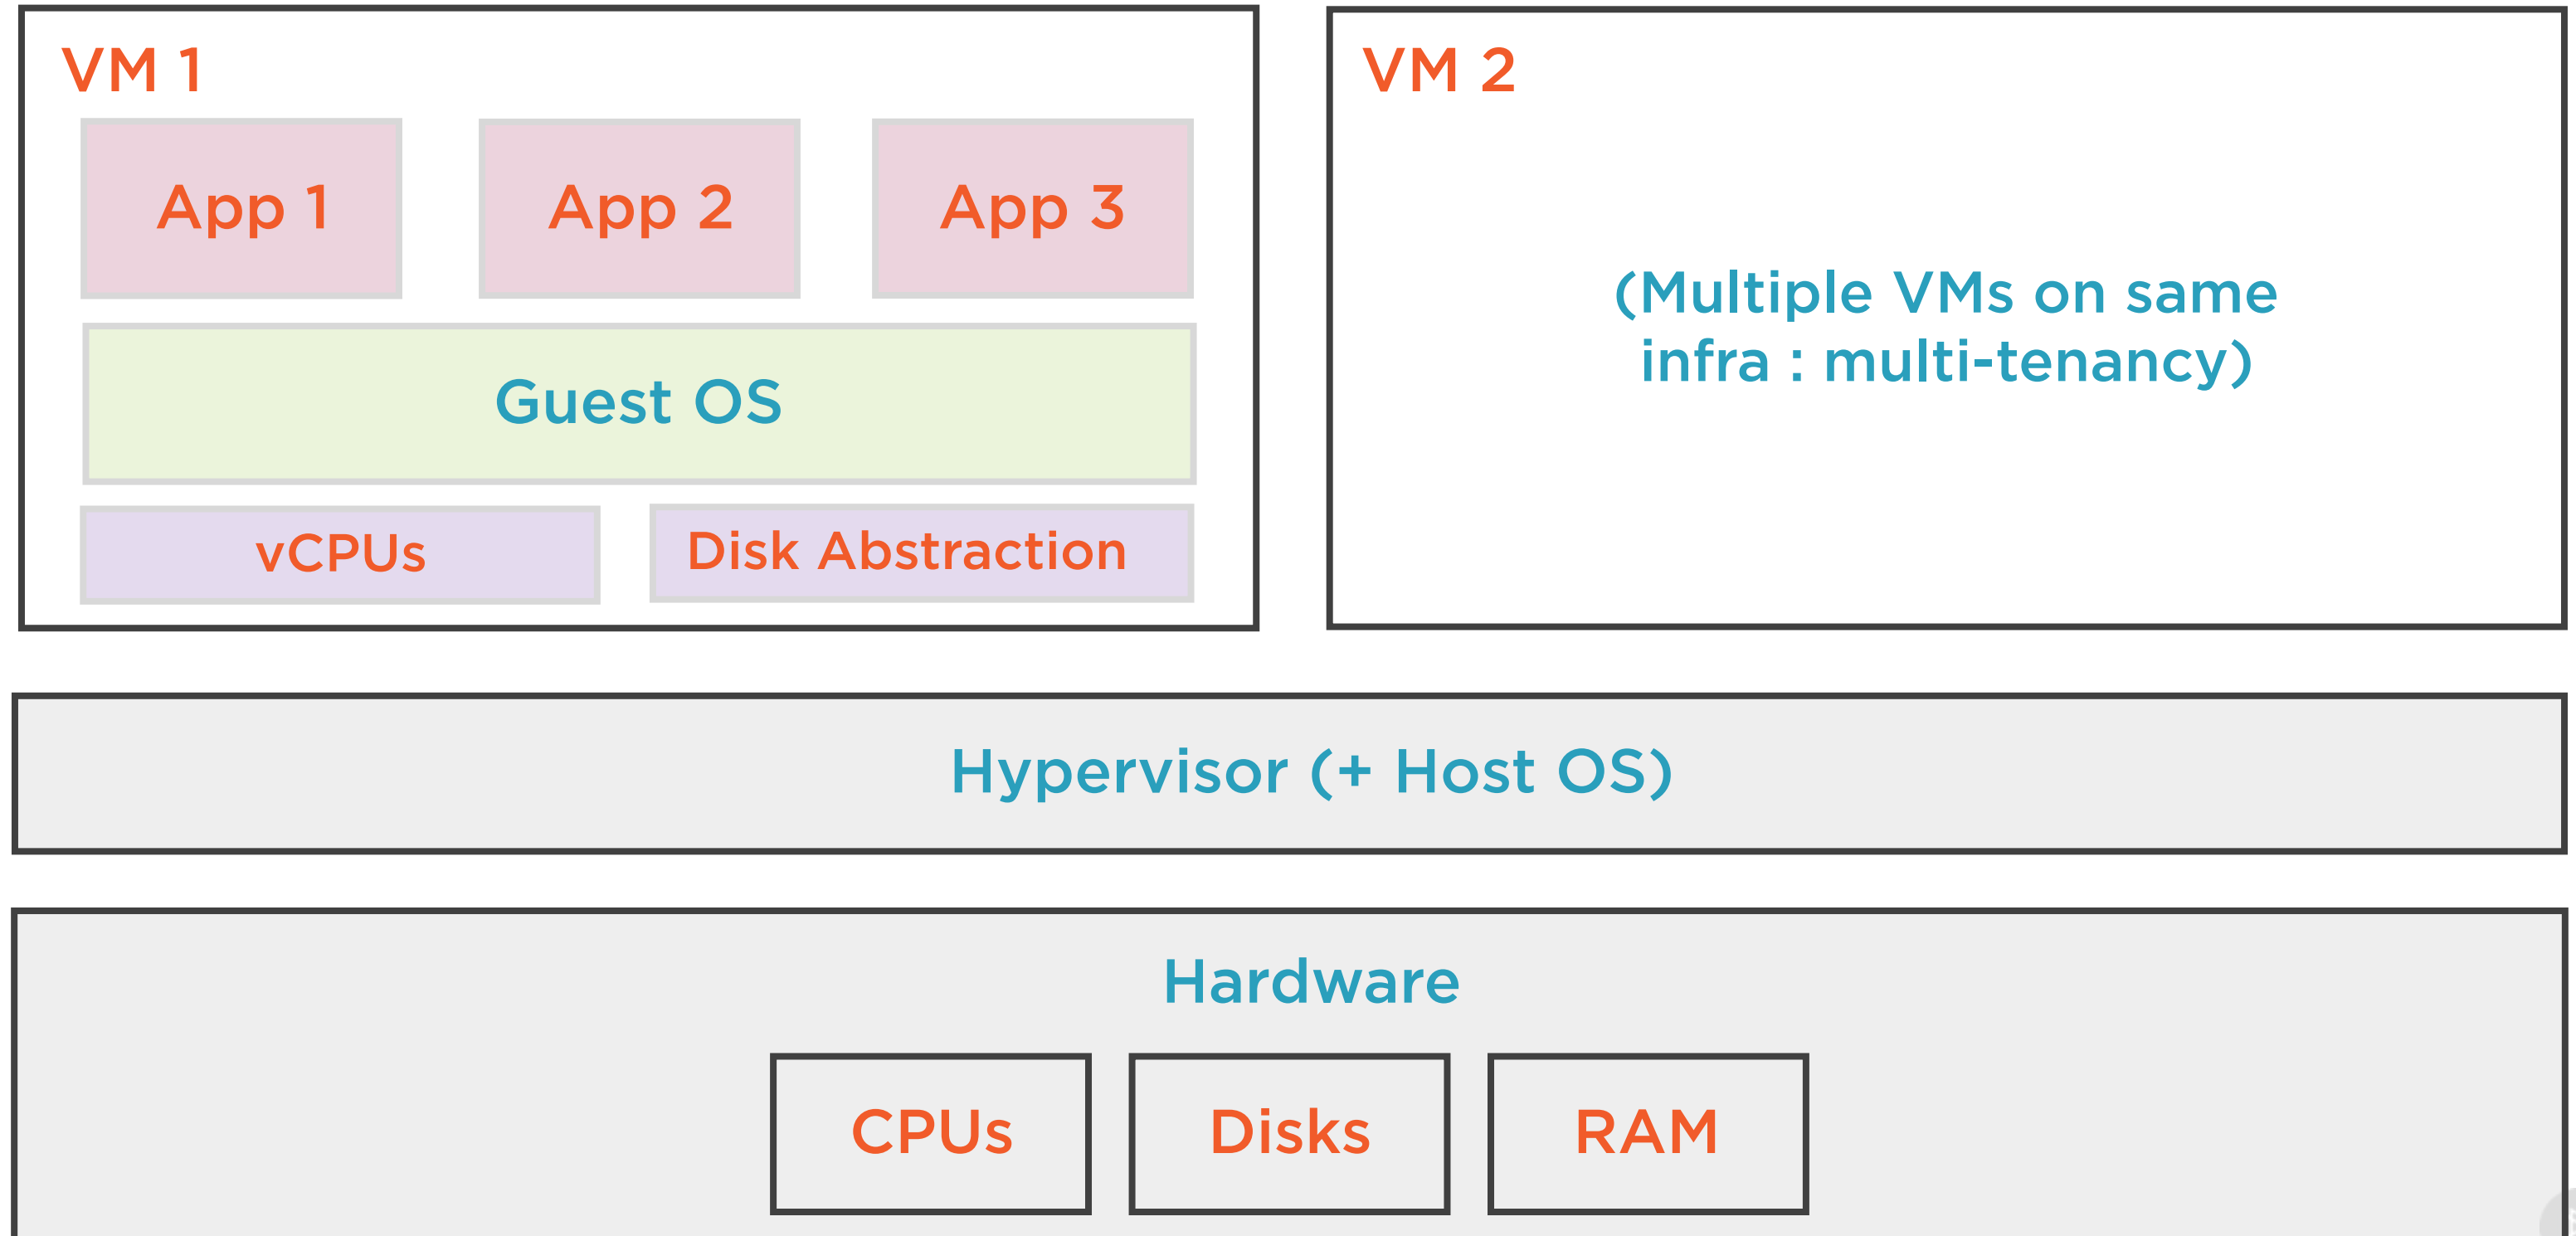
## Bare Metal



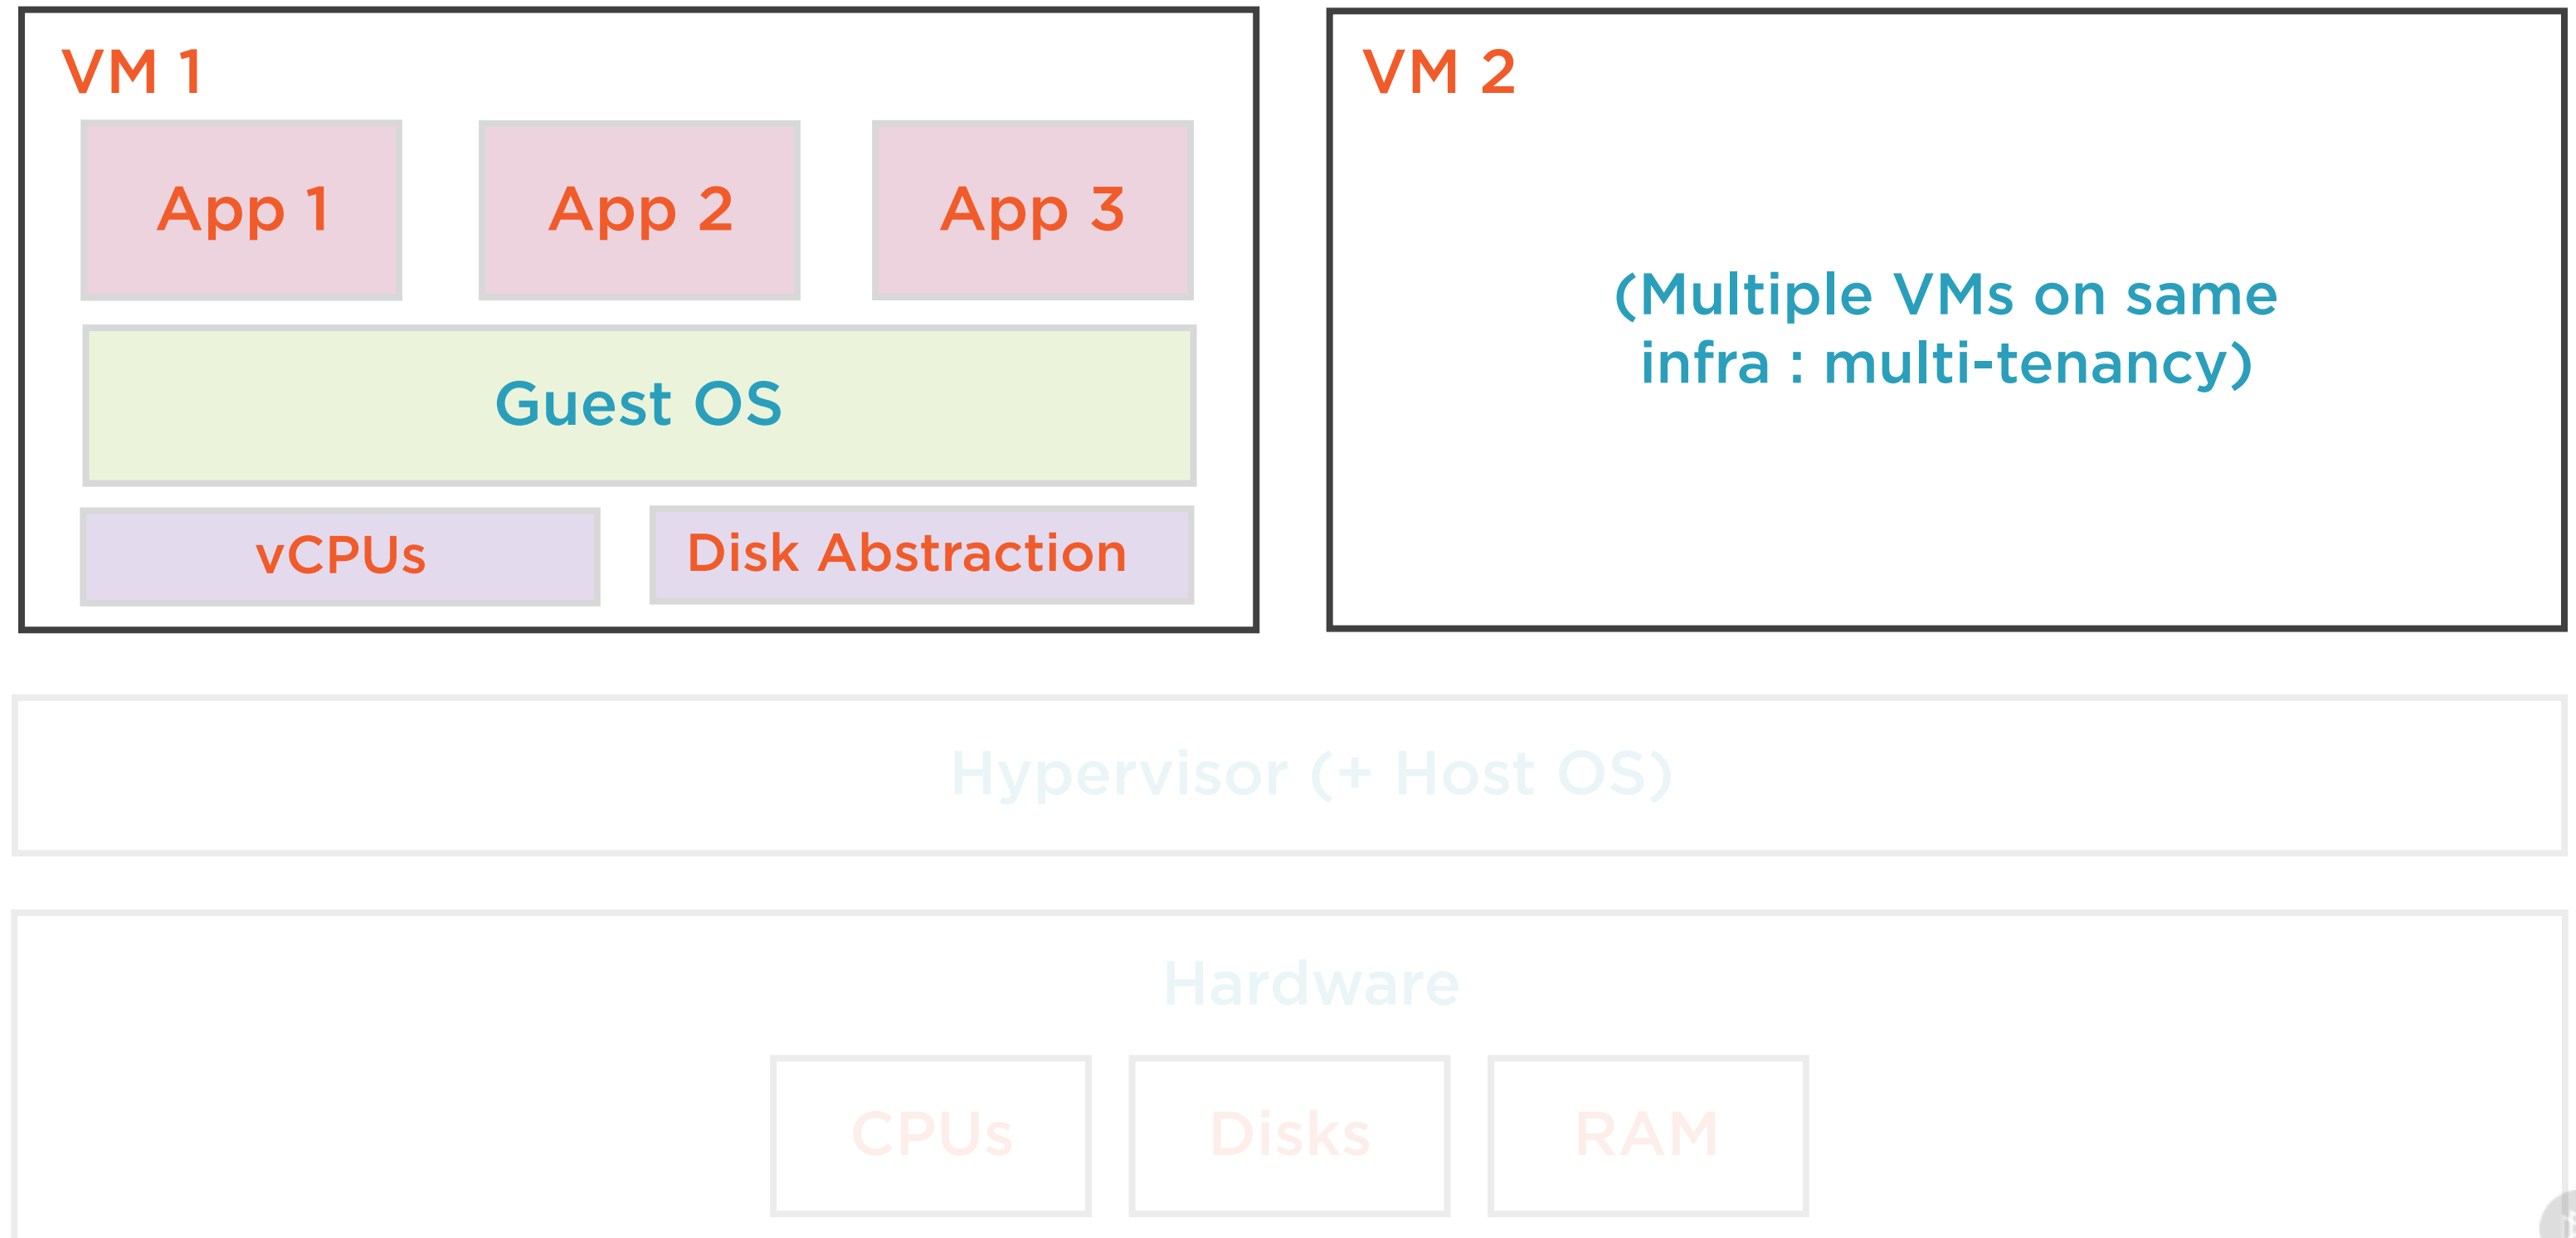
## Virtual Machine



# Apps on Virtual Machines




# Apps on Virtual Machines



VMs are far more lightweight than bare metal environments and are often used to run applications



# Drawbacks of VMs



## **Contain guest OS**

- Introduces platform dependency
- Bloats image size to GB (apps far smaller)

## **Heavyweight**

- Slow to boot up

## **Not trivial to migrate**

- VM migration tools needed

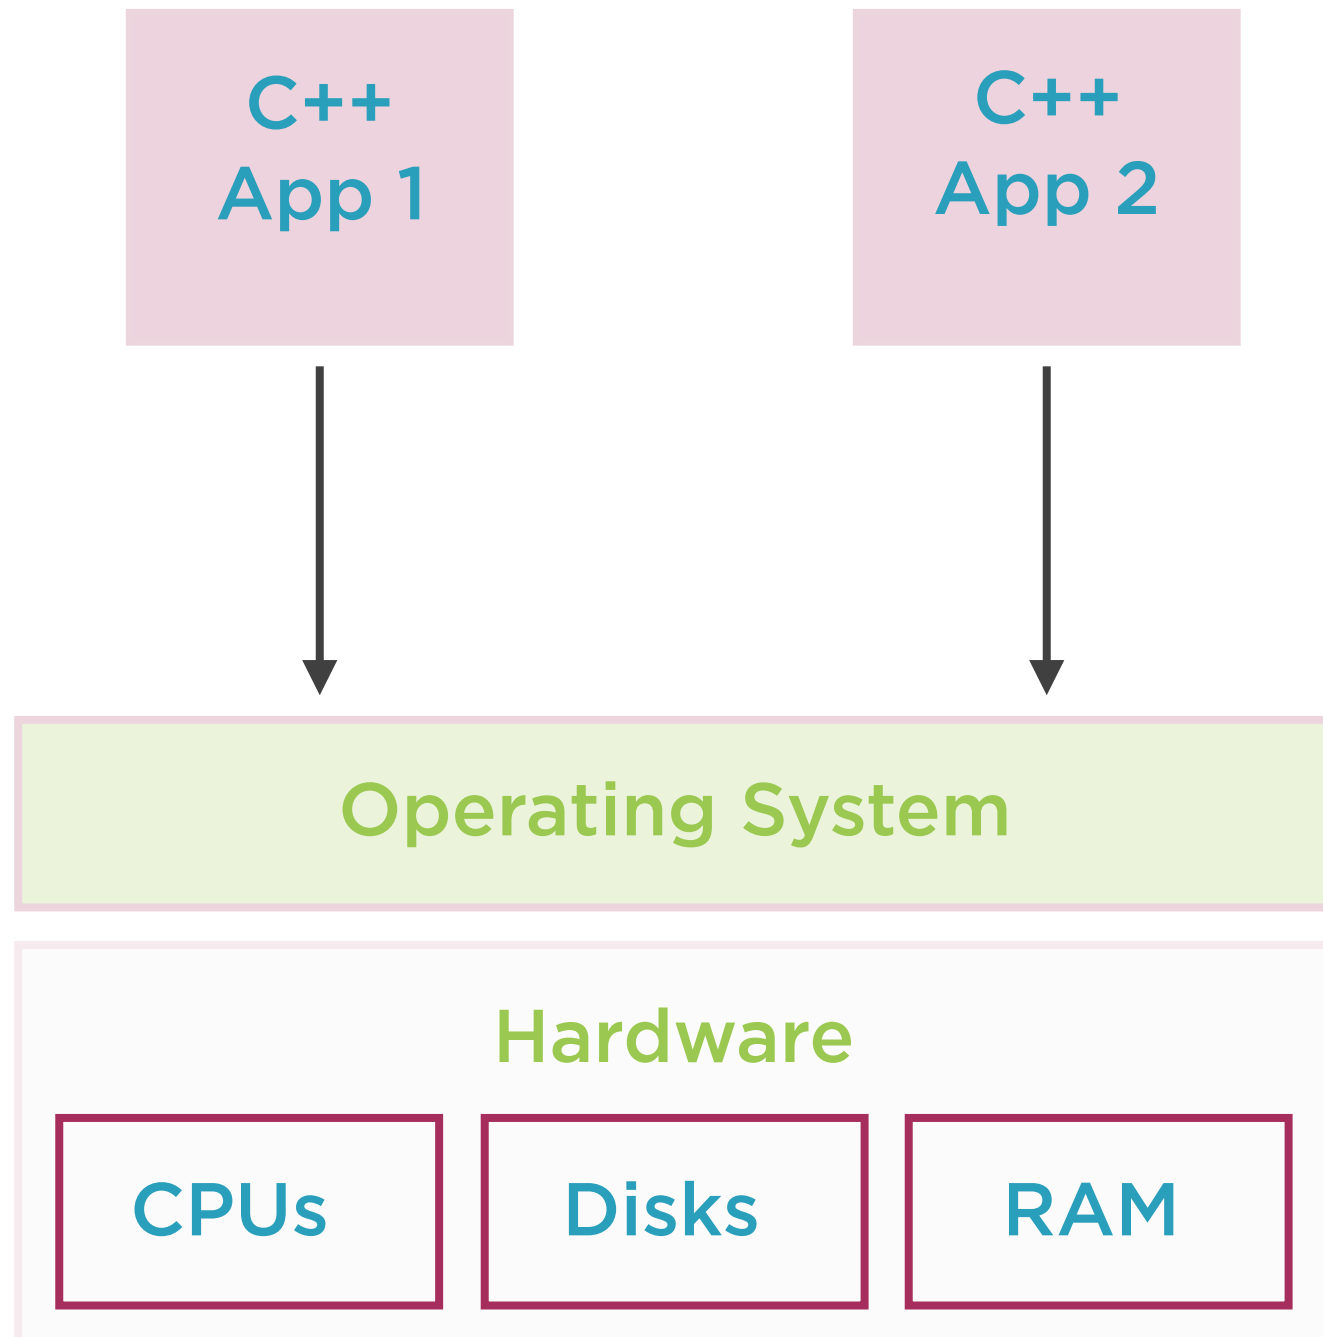
# Container

A container image is a lightweight, stand-alone, executable package of a piece of software that includes everything needed to run it: code, runtime, system tools, system libraries, settings

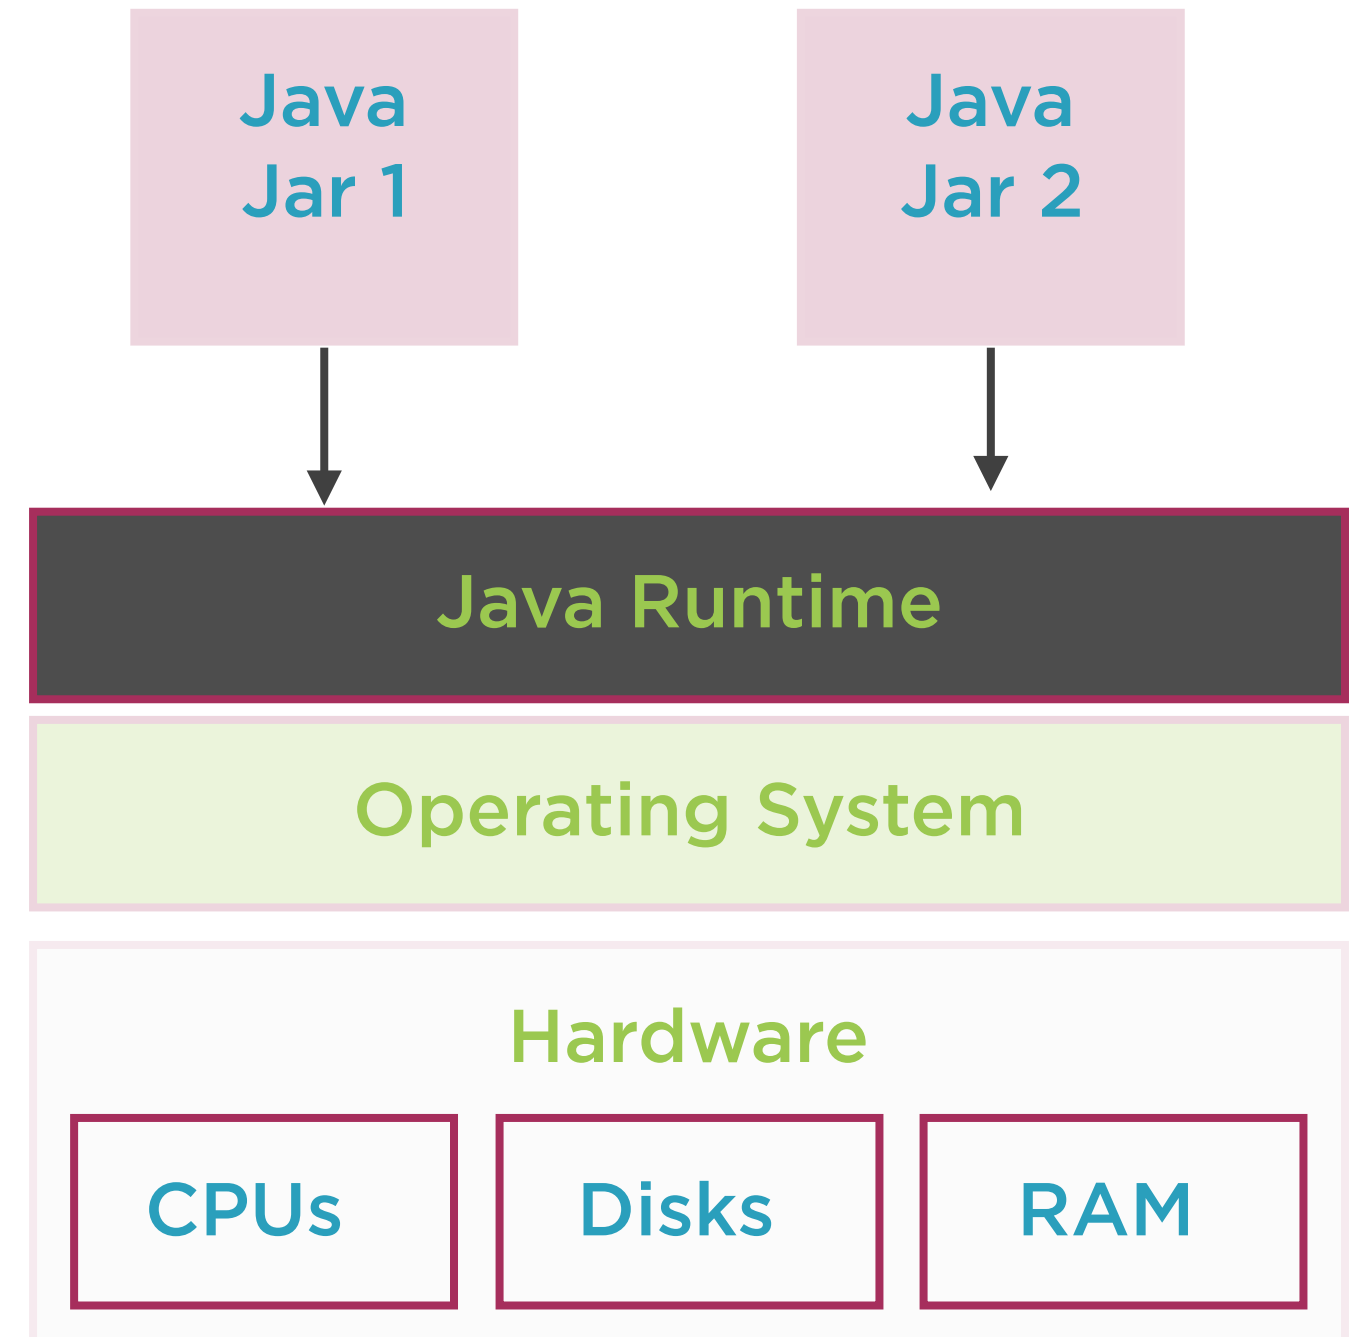


# Containers Are 'Like' Jars

## C++ App

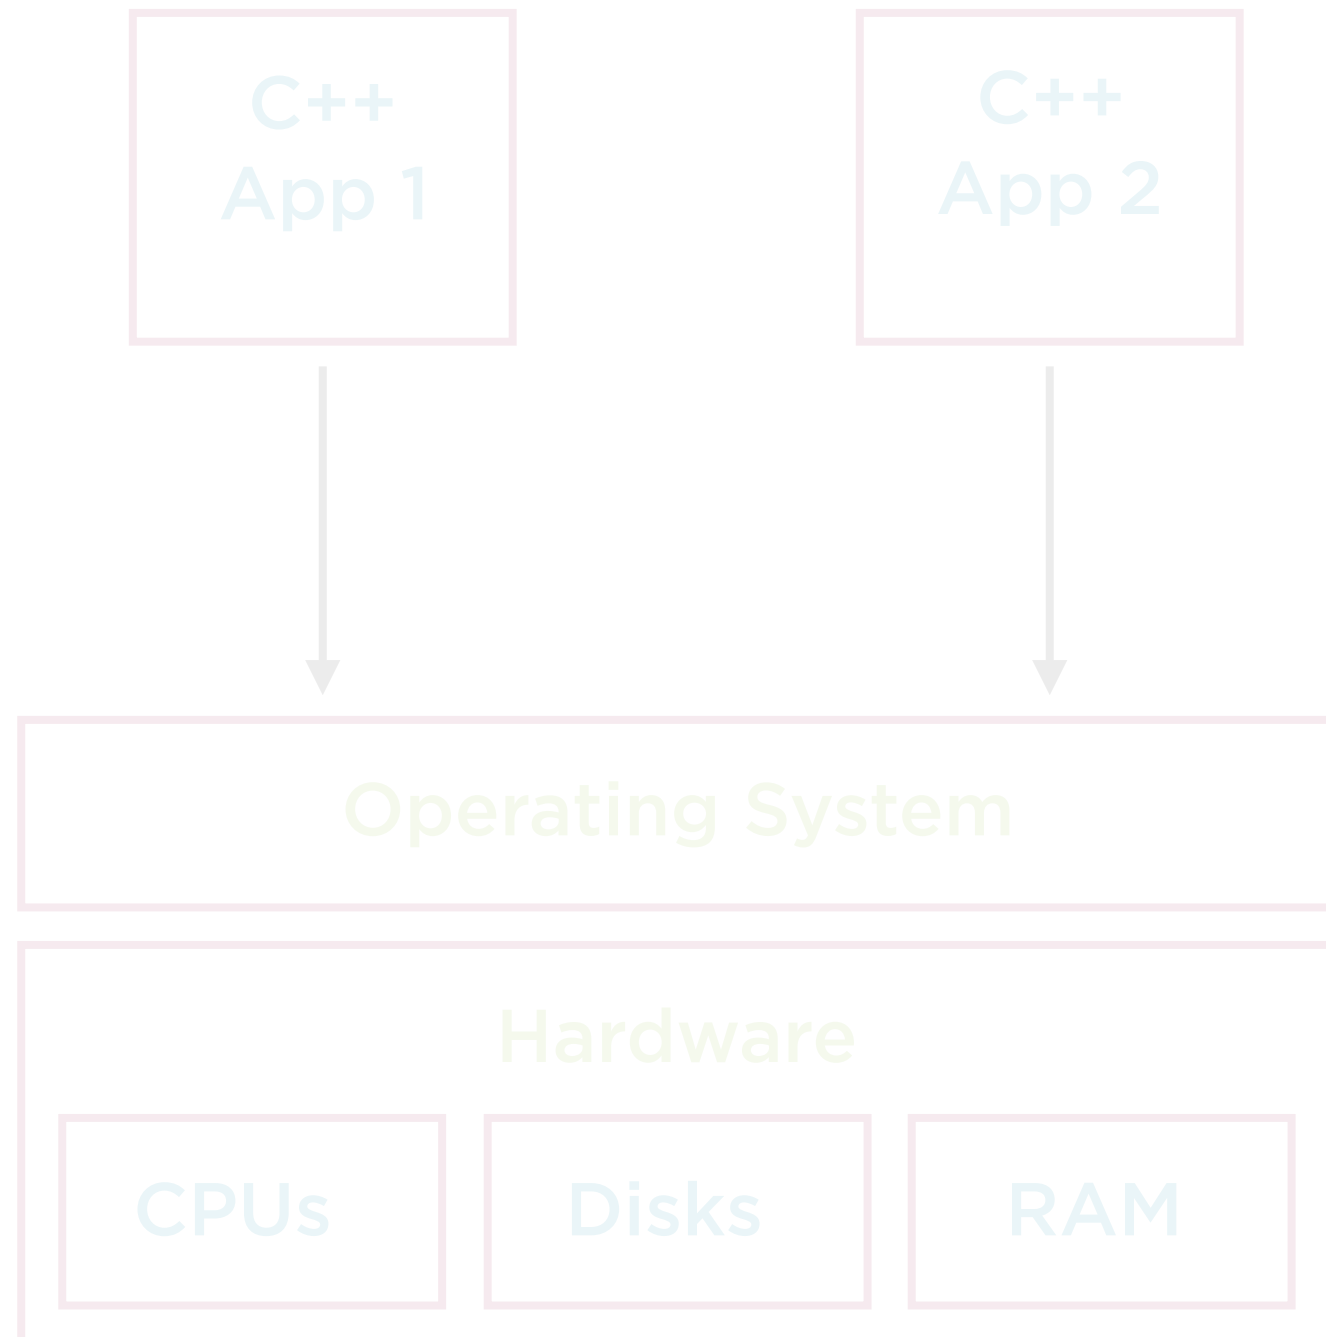


## Java Jars

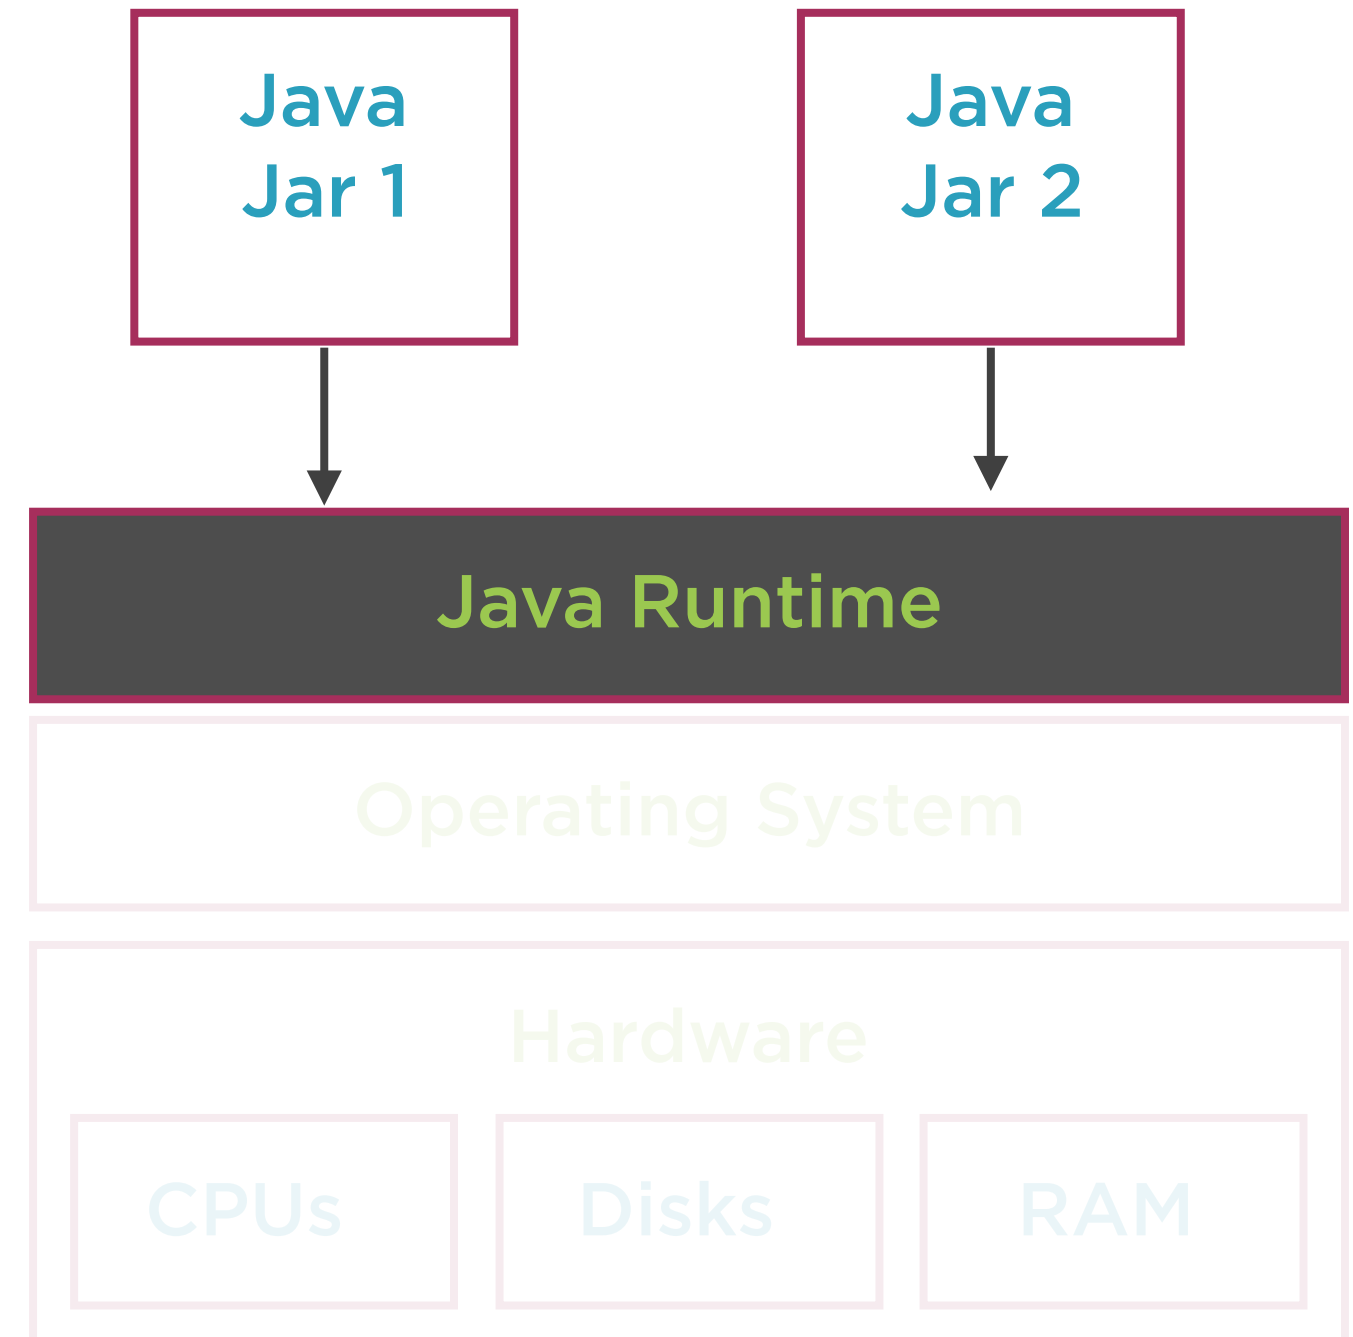


# Containers Are 'Like' Jars

## C++ App



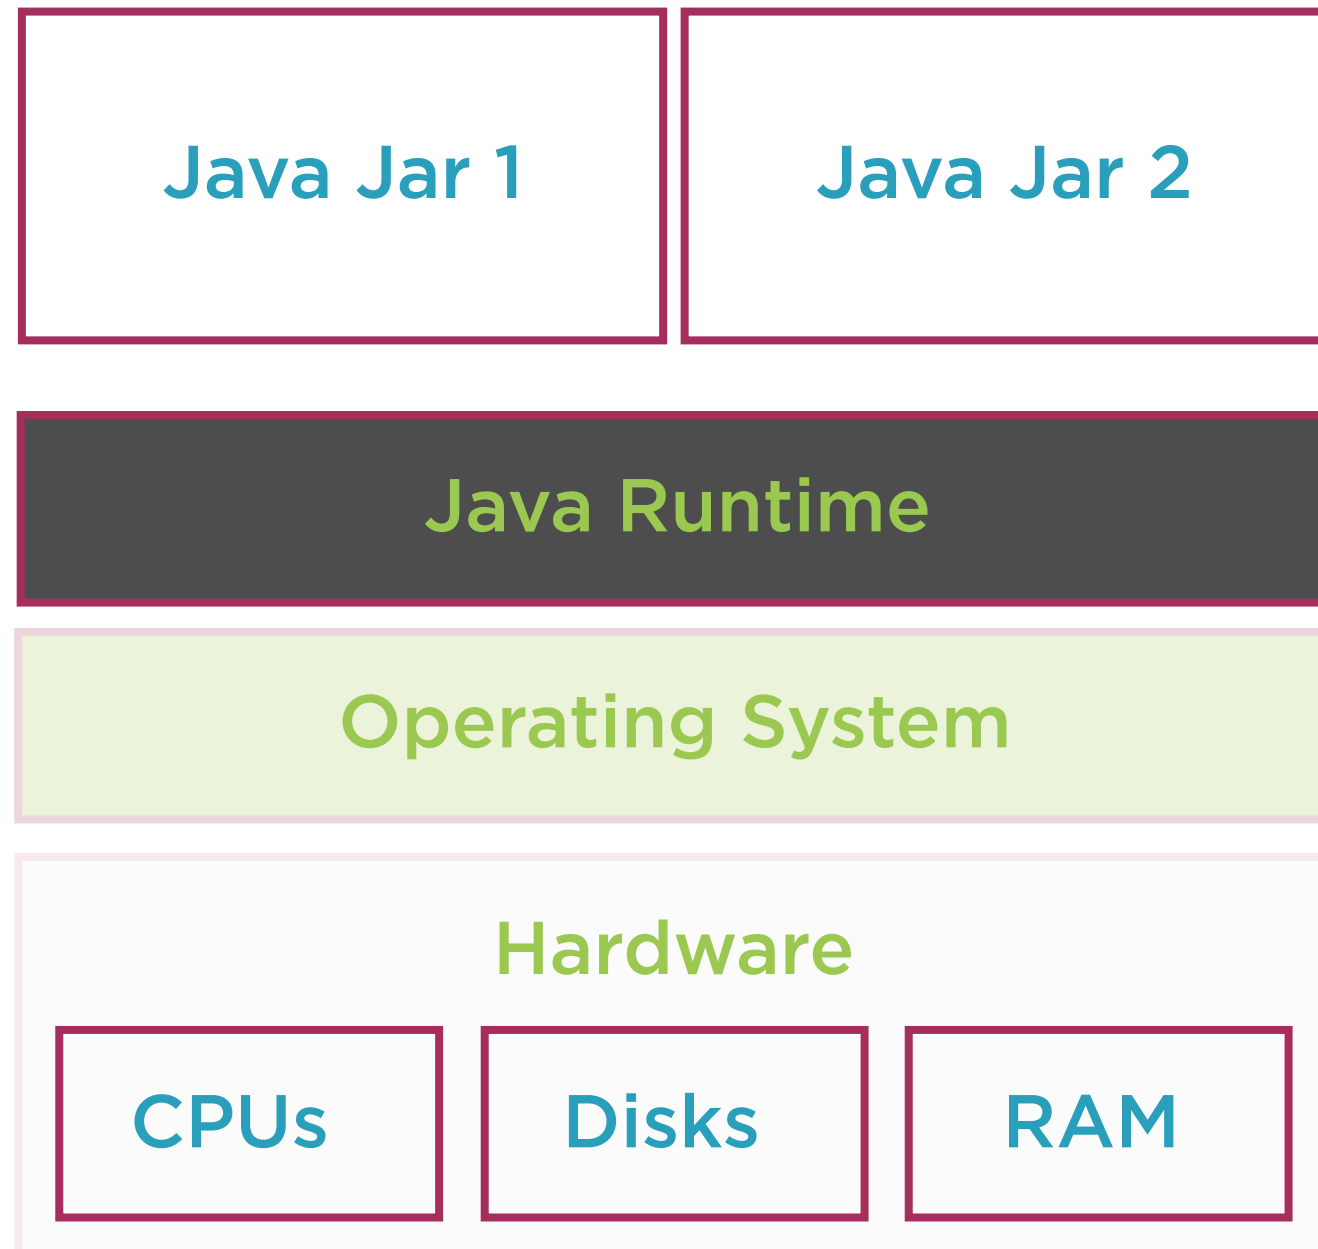
## Java Jars



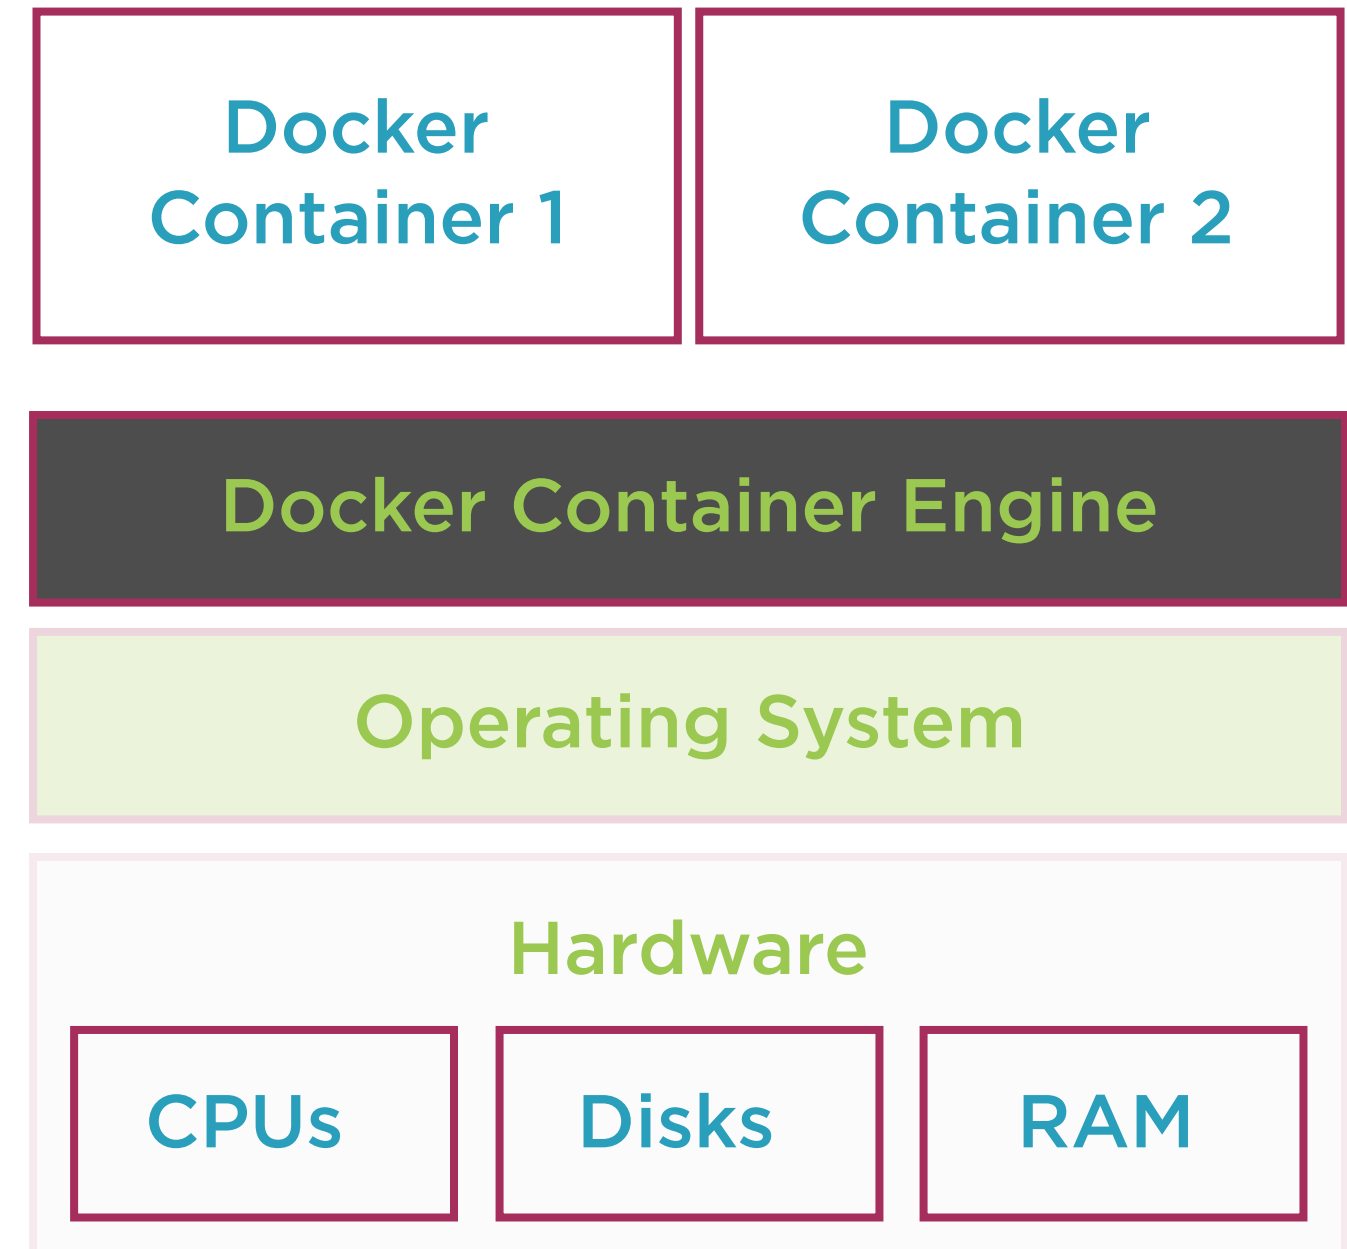


# Containers Are 'Like' Jars

## Java Jars



## Docker Containers



# Java Jars and Containers

## Java Jar Files

Files containing apps

Platform independent

Run on layer of abstraction

Java Runtime

## Containers

Files containing apps

Platform independent

Run on layer of abstraction


Docker Runtime



# Container

A container image is a lightweight, stand-alone, executable package of a piece of software that includes everything needed to run it: code, runtime, system tools, system libraries, settings

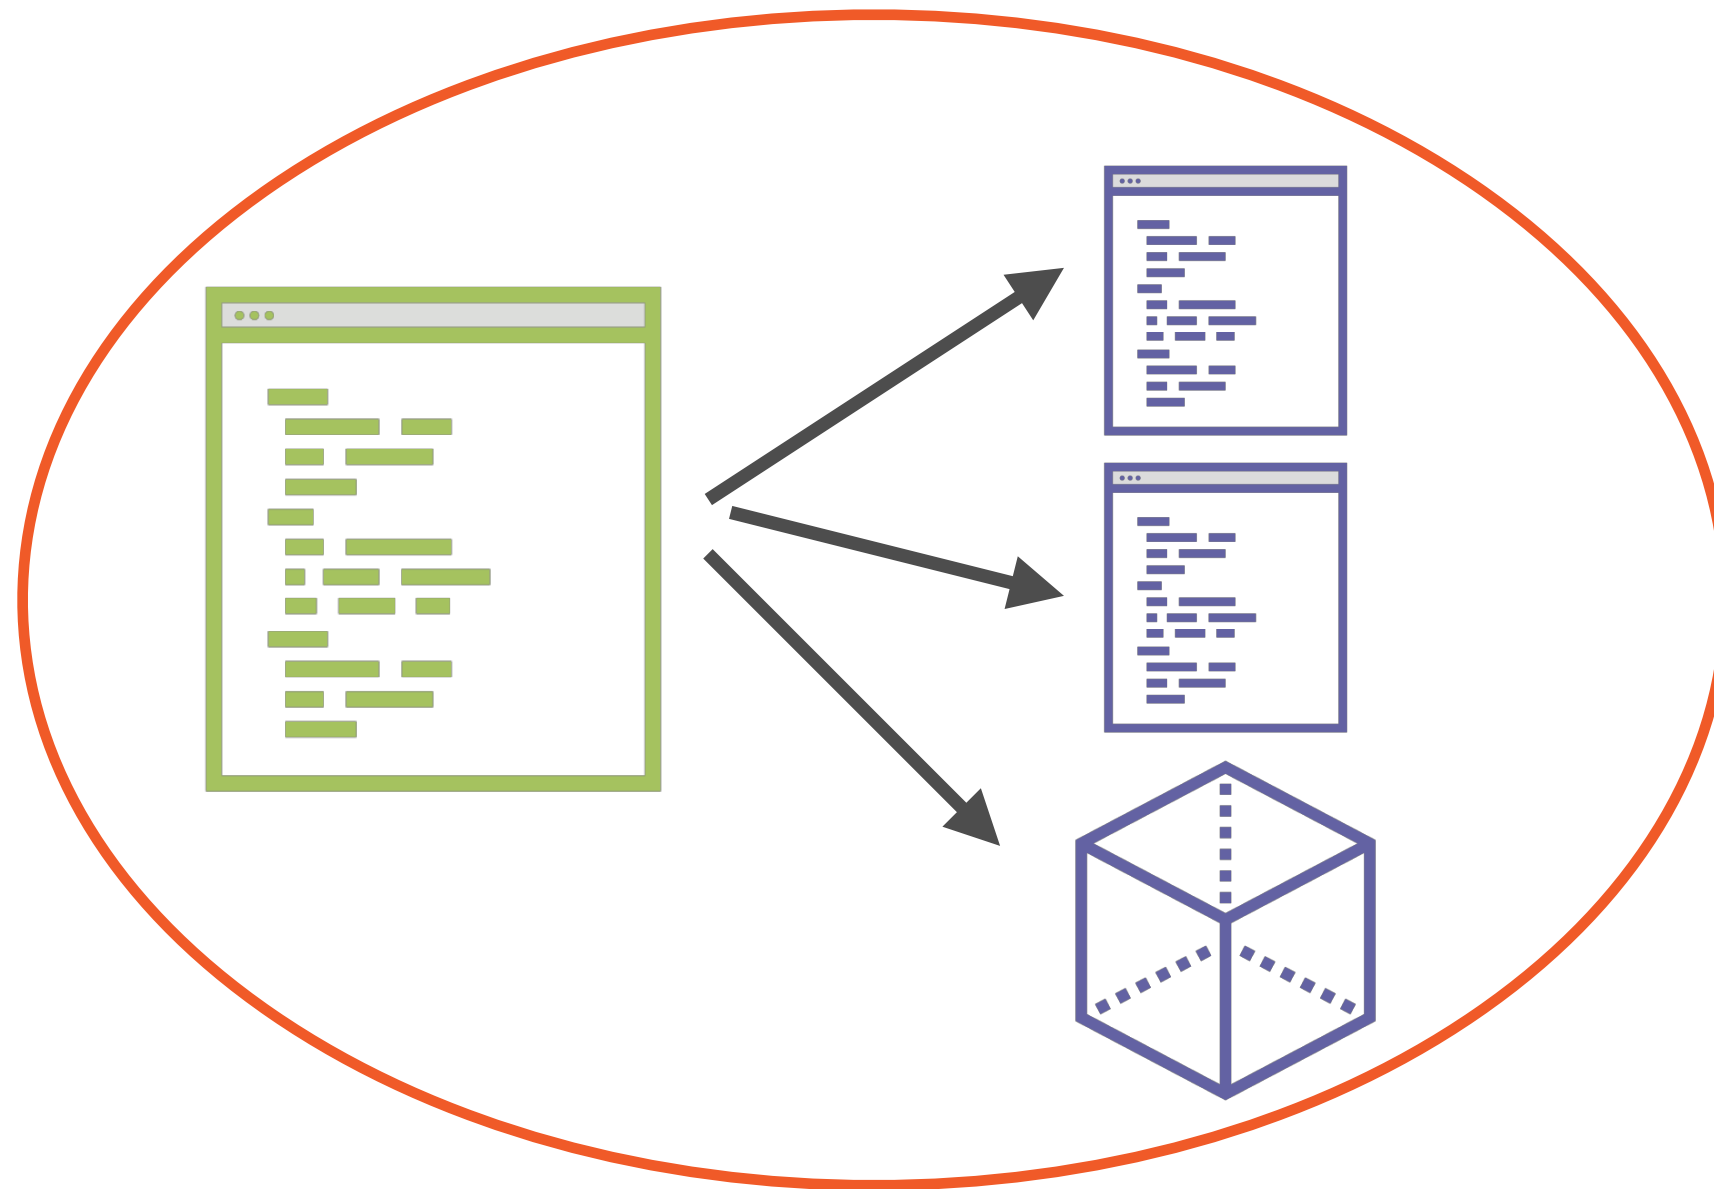
# Docker Containers



A container packages code and its dependencies into an image




# Docker Containers



**A fully self-contained environment**




# Docker Containers



Code is abstracted away from the machine



# Docker Containers



Create as many containers as you want  
from the same image



# Introducing Kubernetes

---





# Attractions of Containers



## No guest OS

- Platform independent
- Considerably smaller than VM images

## Lightweight

- Small and fast
- Quick to start
- Speeds up autoscaling

## Hybrid, multi-cloud

- Hybrid: Work on-premise and on cloud
- Multi-cloud: Not tied to any specific cloud platform



# Standalone Container Limitations



## No autohealing

- Crashed containers won't restart automatically
- Need higher level orchestration

## No scaling or autoscaling

- Overloaded containers don't spawn more automatically
- Need higher level orchestration

## No load balancing

- Containers can't share load automatically
- Need higher level orchestration

## No isolation

- Crashing containers can take each other down
- Need sandbox to separate them



# Kubernetes

Orchestration technology for containers - convert isolated containers running on different hardware into a cluster



# Compute Choices



Less ops (admin, provisioning)

More control, low-level access



# Compute Choices



Less ops (admin, provisioning)

More control, low-level access



Kubernetes is fast emerging as middle-ground between IaaS and PaaS in a hybrid, multi-cloud world





# Hybrid, Multi-cloud

## **Hybrid: Runs on-premise and on cloud**

- Provides smooth migration path

## **Multi-cloud: Supported by all big cloud platforms**

- Important for strategic independence
- Amazon-Whole Foods merger



# Compute Choices



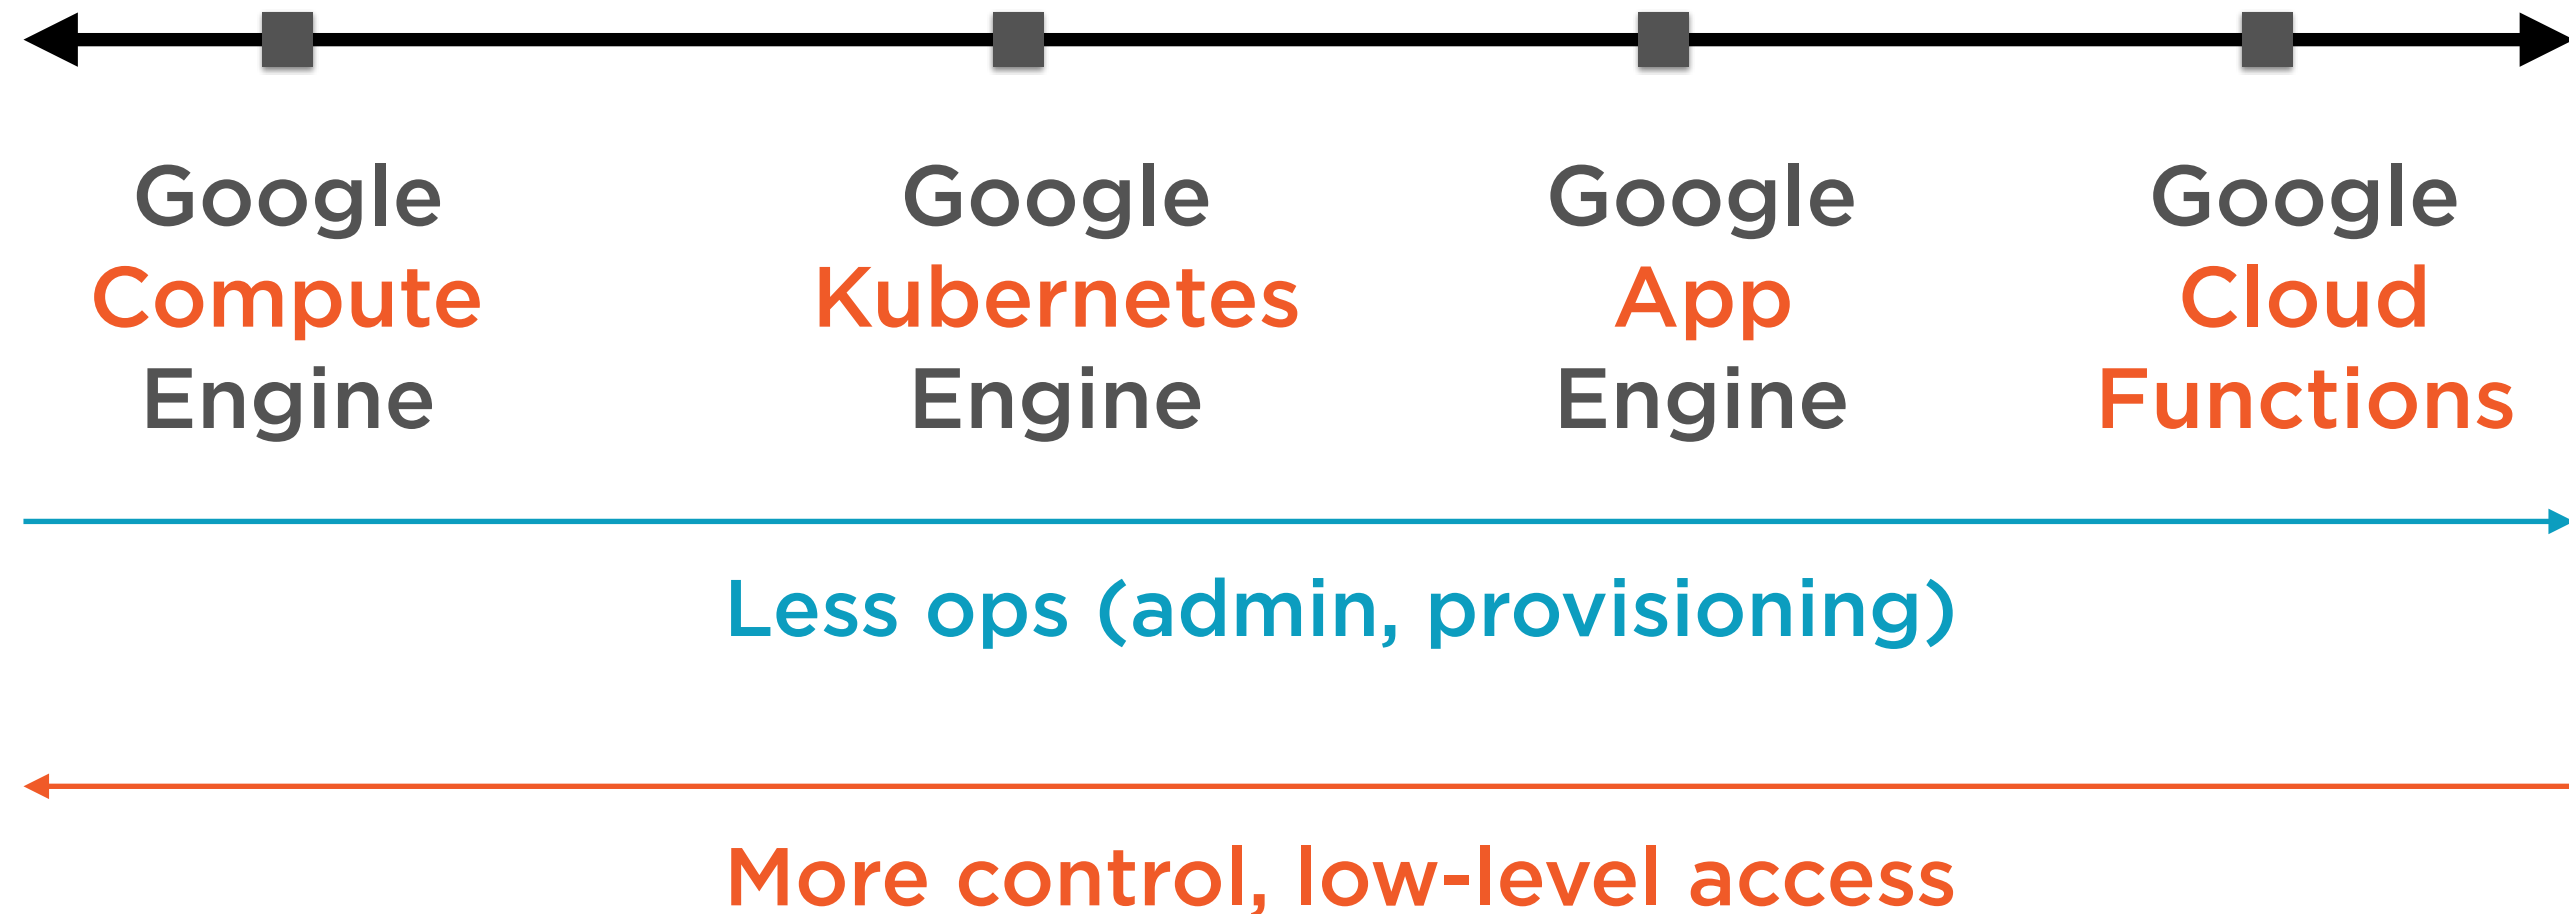
Less ops (admin, provisioning)

More control, low-level access

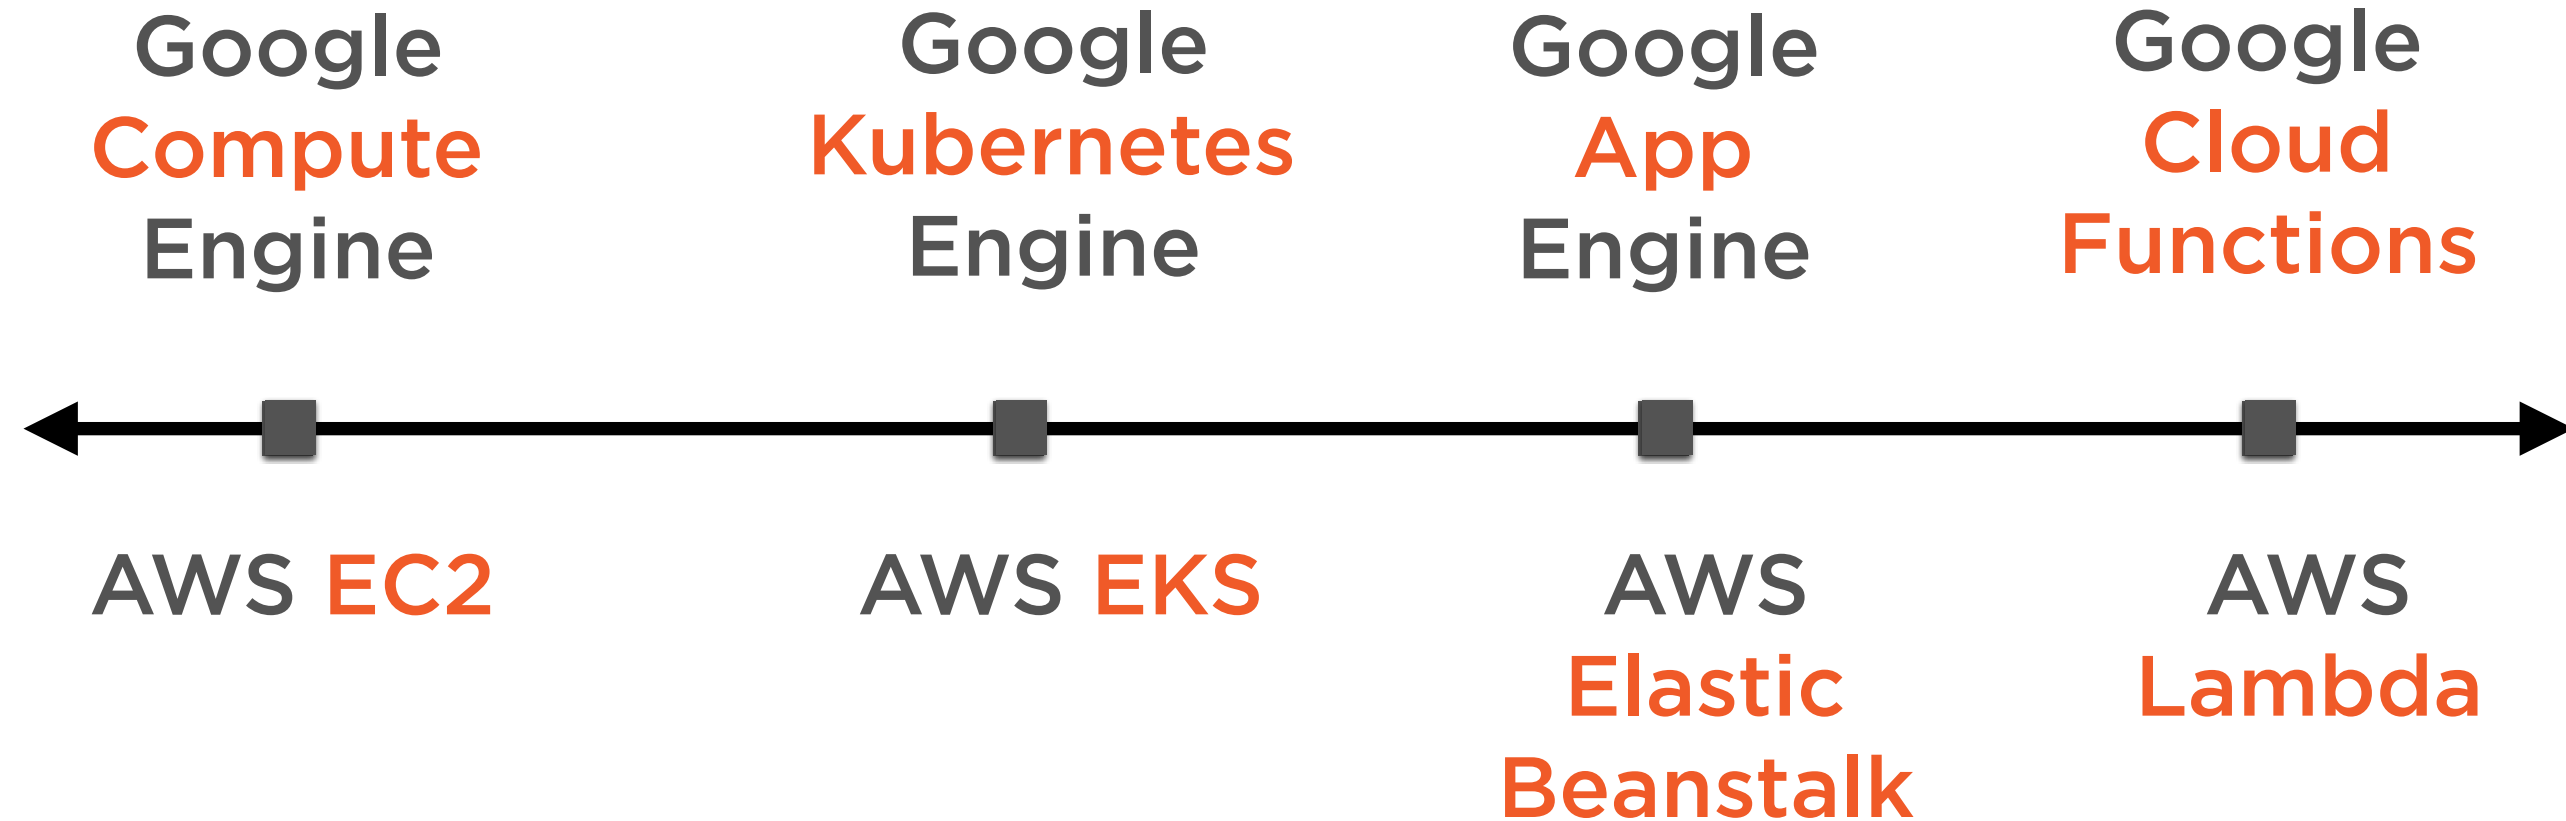




# GCP Compute Choices



# GCP Compute Choices



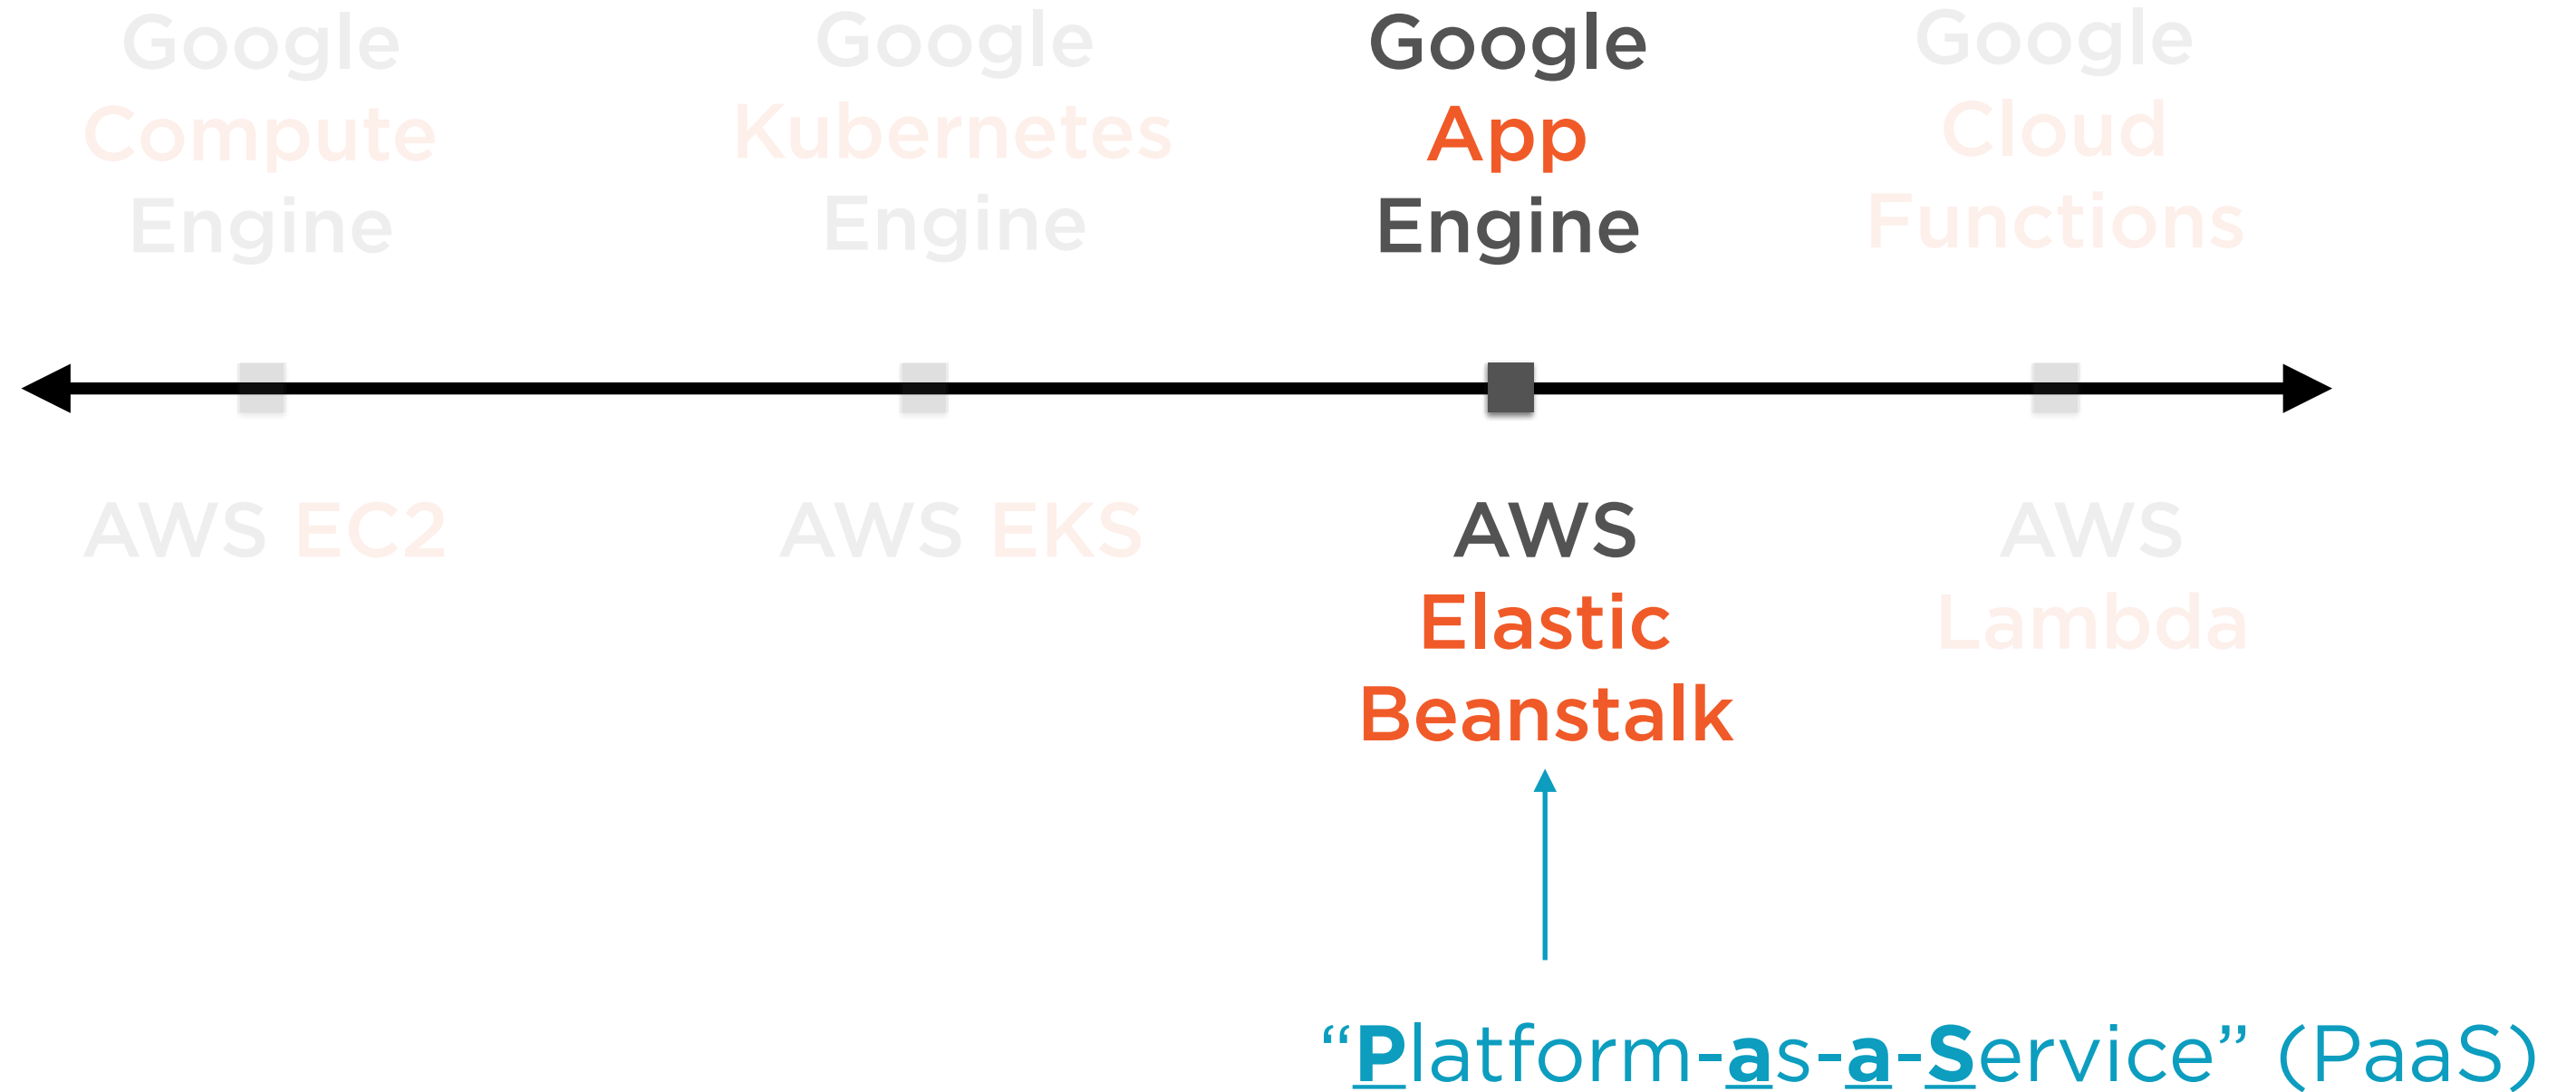
Every major cloud platform supports the same range of compute choices



# GCP Compute Choices



# GCP Compute Choices



# IaaS vs. PaaS

## Infrastructure-as-a-Service

Heavy operational burden

Migration is hard

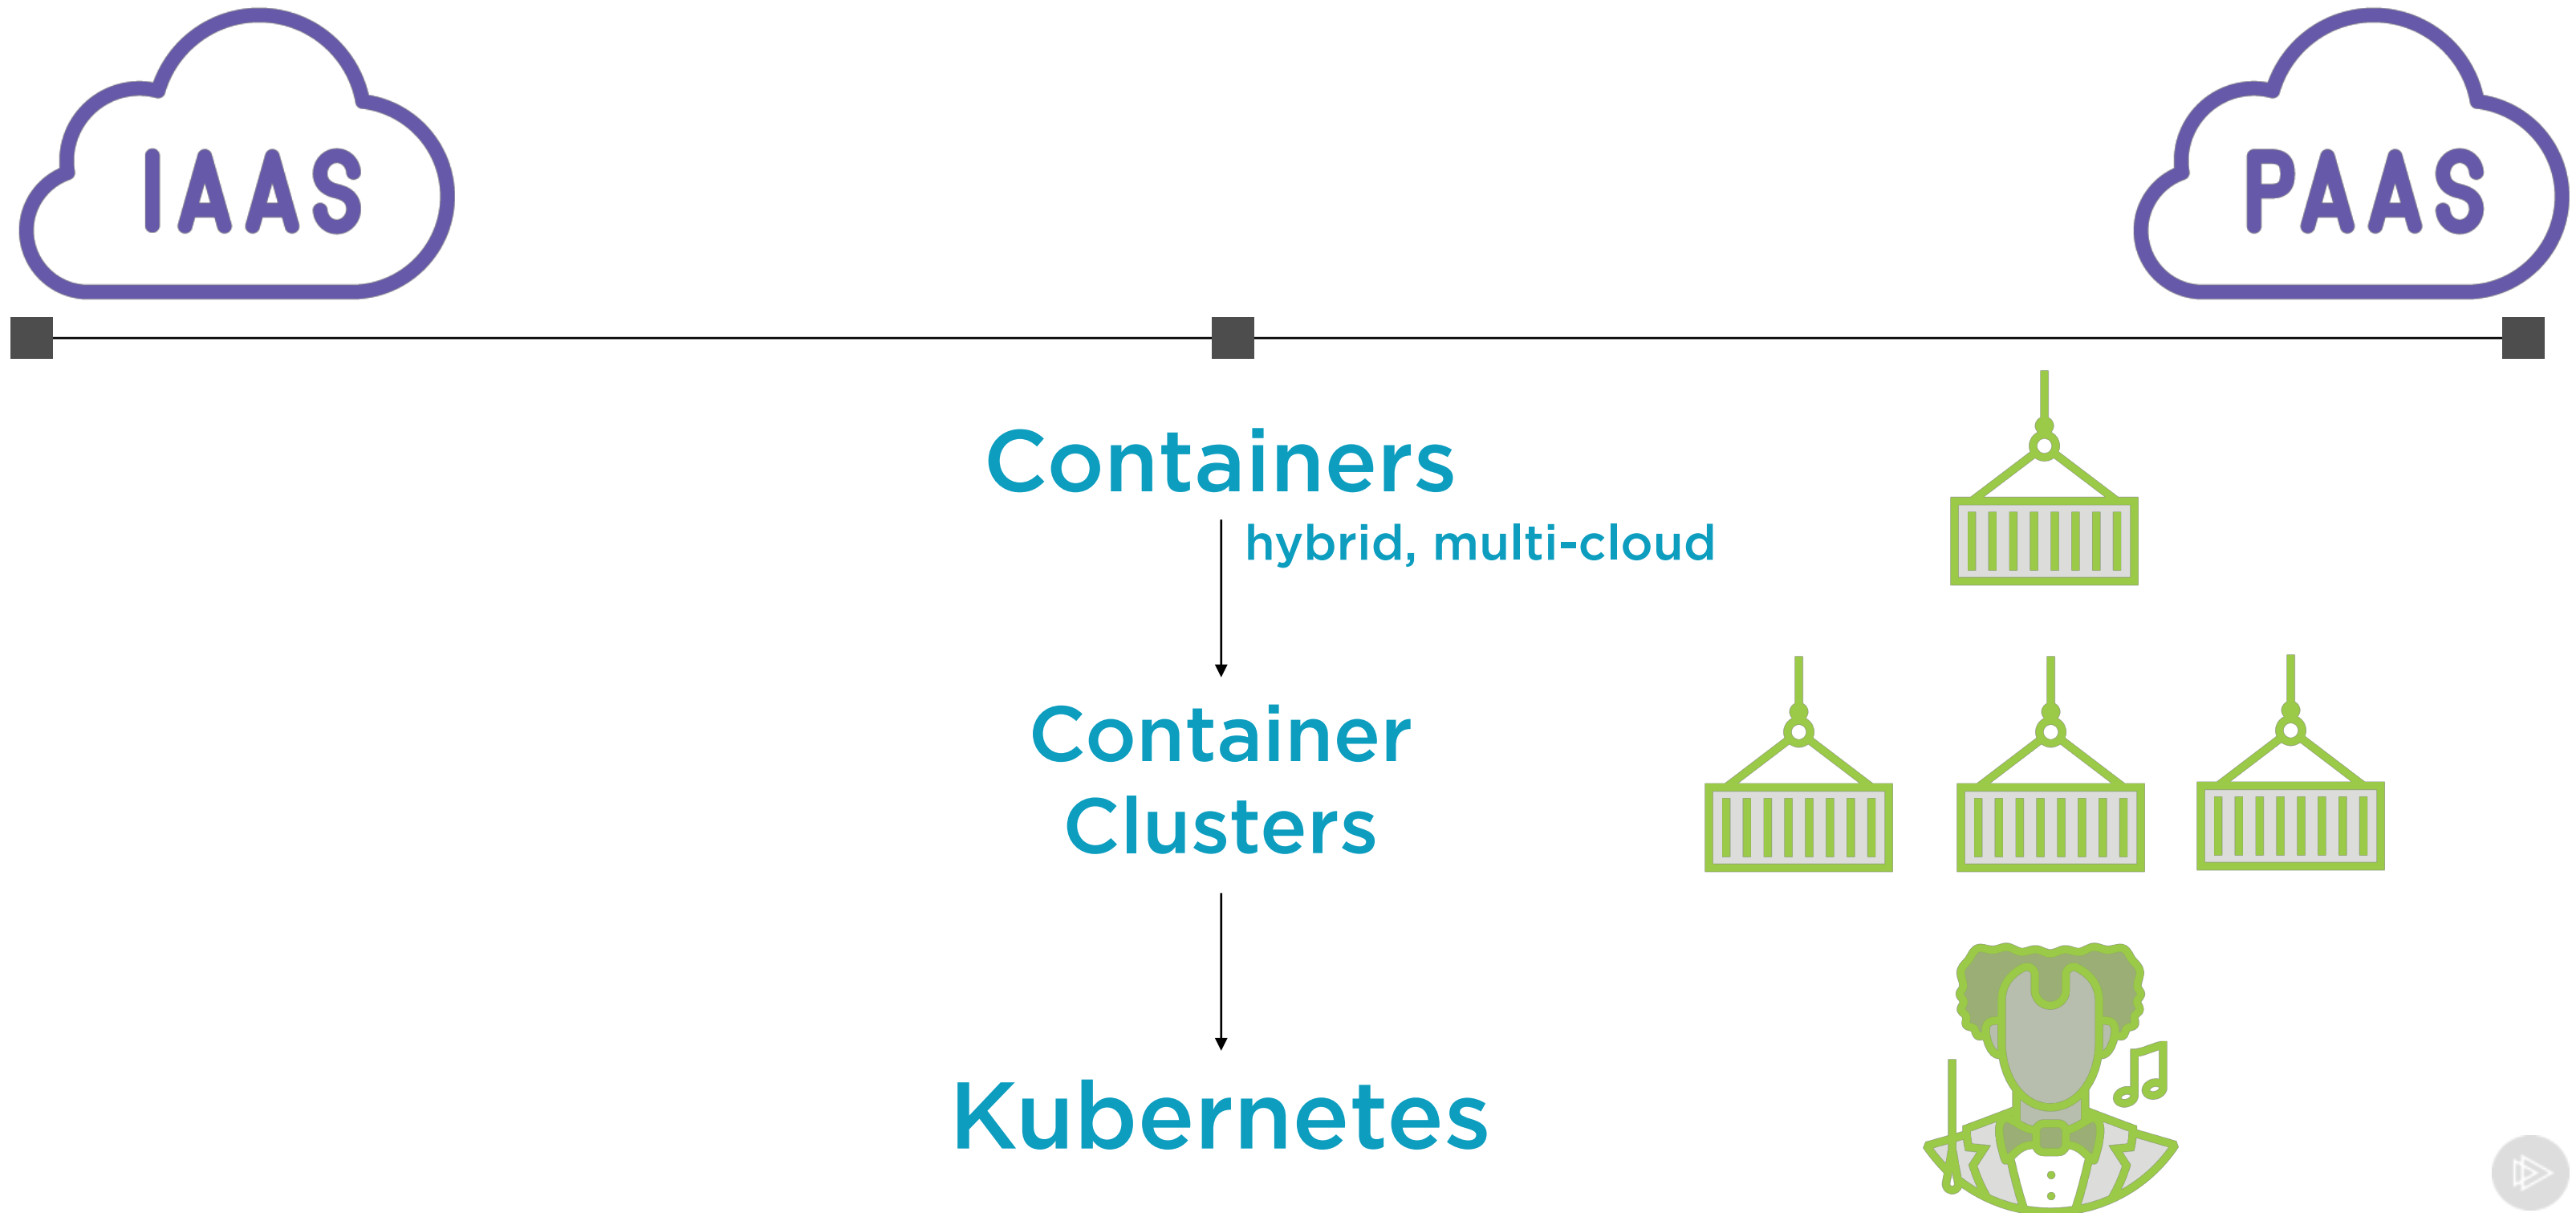
## Platform-as-a-Service

Provider lock-in

Migration is very hard



# Compute Choices



# Kubernetes as Orchestrator



**Fault-tolerance**

**Autohealing**

**Isolation**

**Scaling**

**Autoscaling**

**Load balancing**



All of these are possible in a  
Kubernetes cluster using higher  
level abstractions







# Hybrid, Multi-cloud

**Kubernetes is supported by each of the Big Three**

**Special relationship with GCP**

**Kubernetes originated at Google**



# Google Kubernetes Engine (GKE)



**Google Kubernetes Engine**

**Service for working with Kubernetes clusters on GCP**


**Runs Kubernetes on GCE VM instances**







# Clusters, Nodes, Node Pools, Node Images

---

# Kubernetes Clusters



Master node



Worker  
nodes





# Master

**One or more nodes designated master nodes**

**Managed by GKE**

**Not visible directly to user**

**Multi-master for high-availability**



**Kubernetes Control Plane directed from here**








All user interactions with Kubernetes clusters are via the **kube-apiserver** running on the master node



# Kubernetes Clusters



# Nodes




Nodes are on-premises or cloud VMs  
on which containers are run











# Node Pools



A subset of node instances which have the same configuration are called node pools




# Node Images



Special operating system images  
are available on the Google Cloud  
to run on Kubernetes nodes





# Node Images

## Container-optimized OS

- Enhances node security
- Supported by teams at Google

## Container-optimized OS with **containerd**

- **containerd** as the main container runtime

## Ubuntu

- Optimized for GKE
- Additional support for XFS, CephFS, Sysdig or Debian packages



# Pods

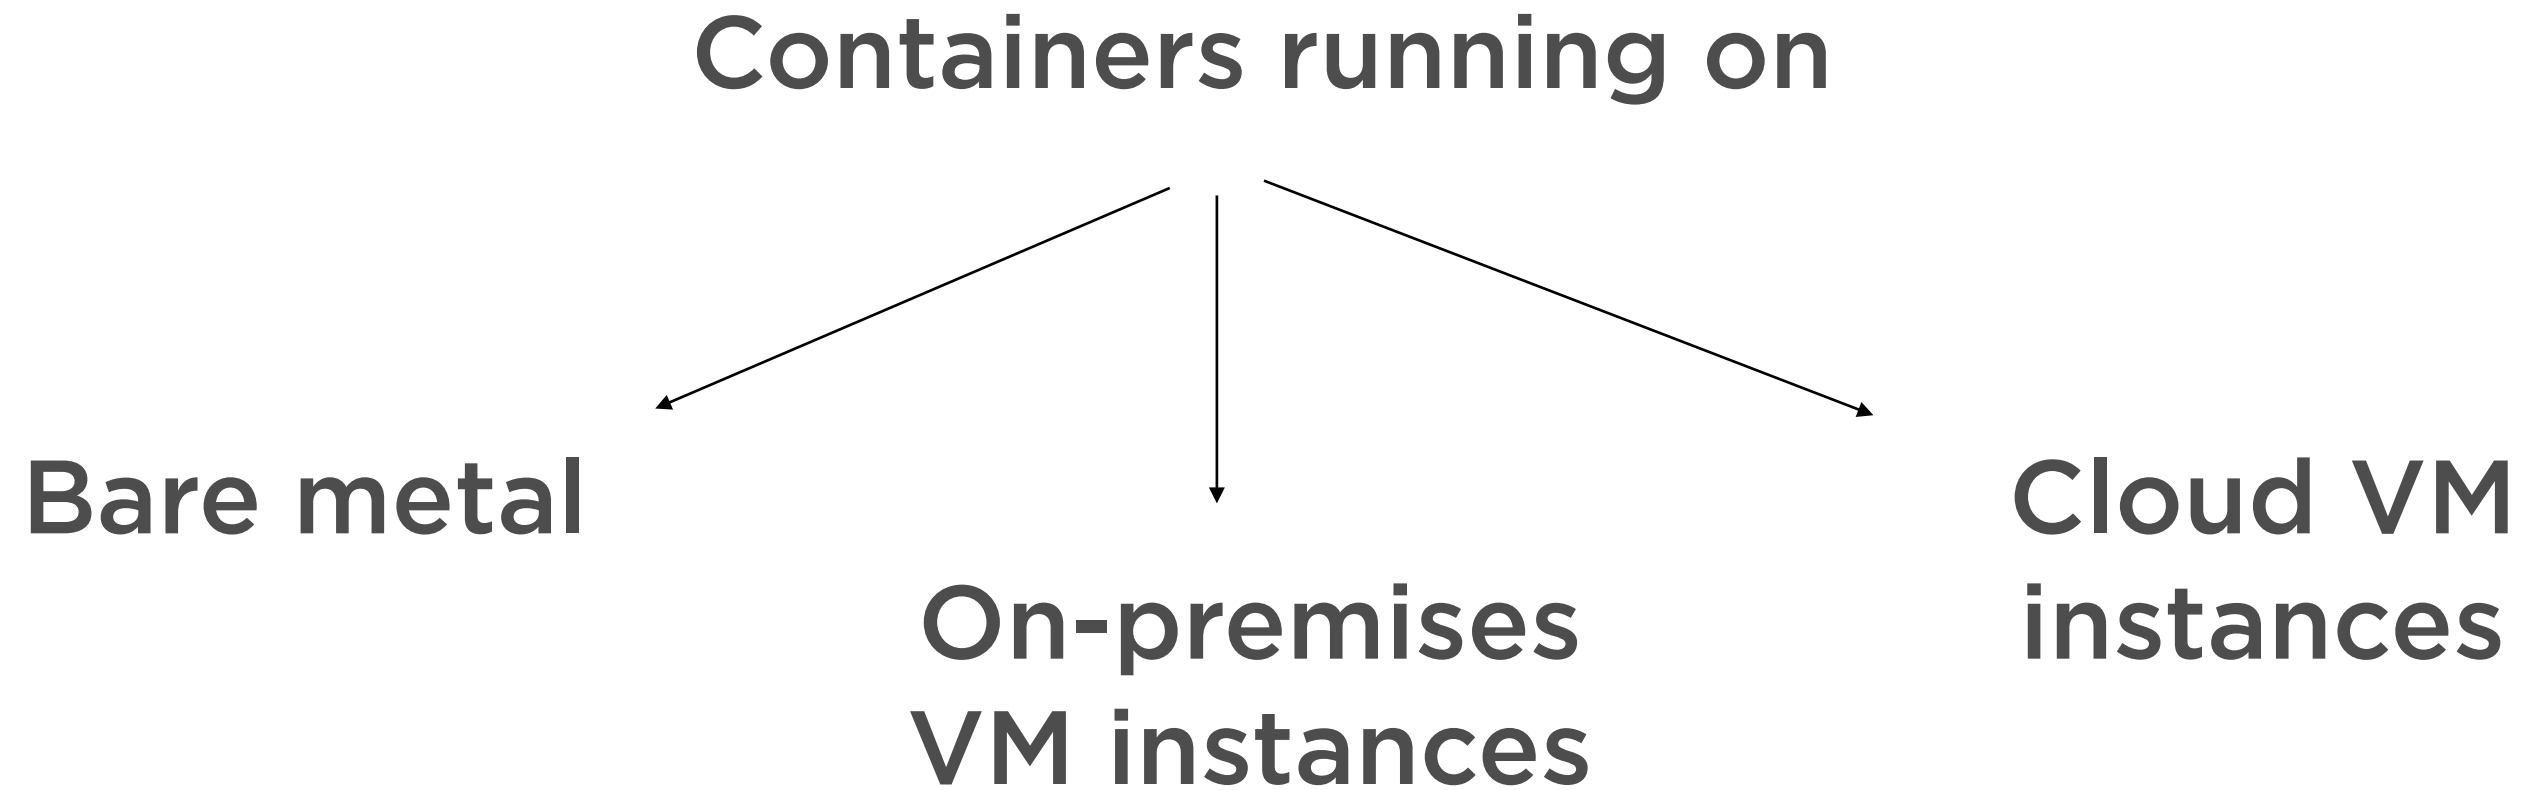
---



# Kubernetes

Orchestration technology for containers - convert isolated containers running on different hardware into a cluster


# Isolated Container Deployments



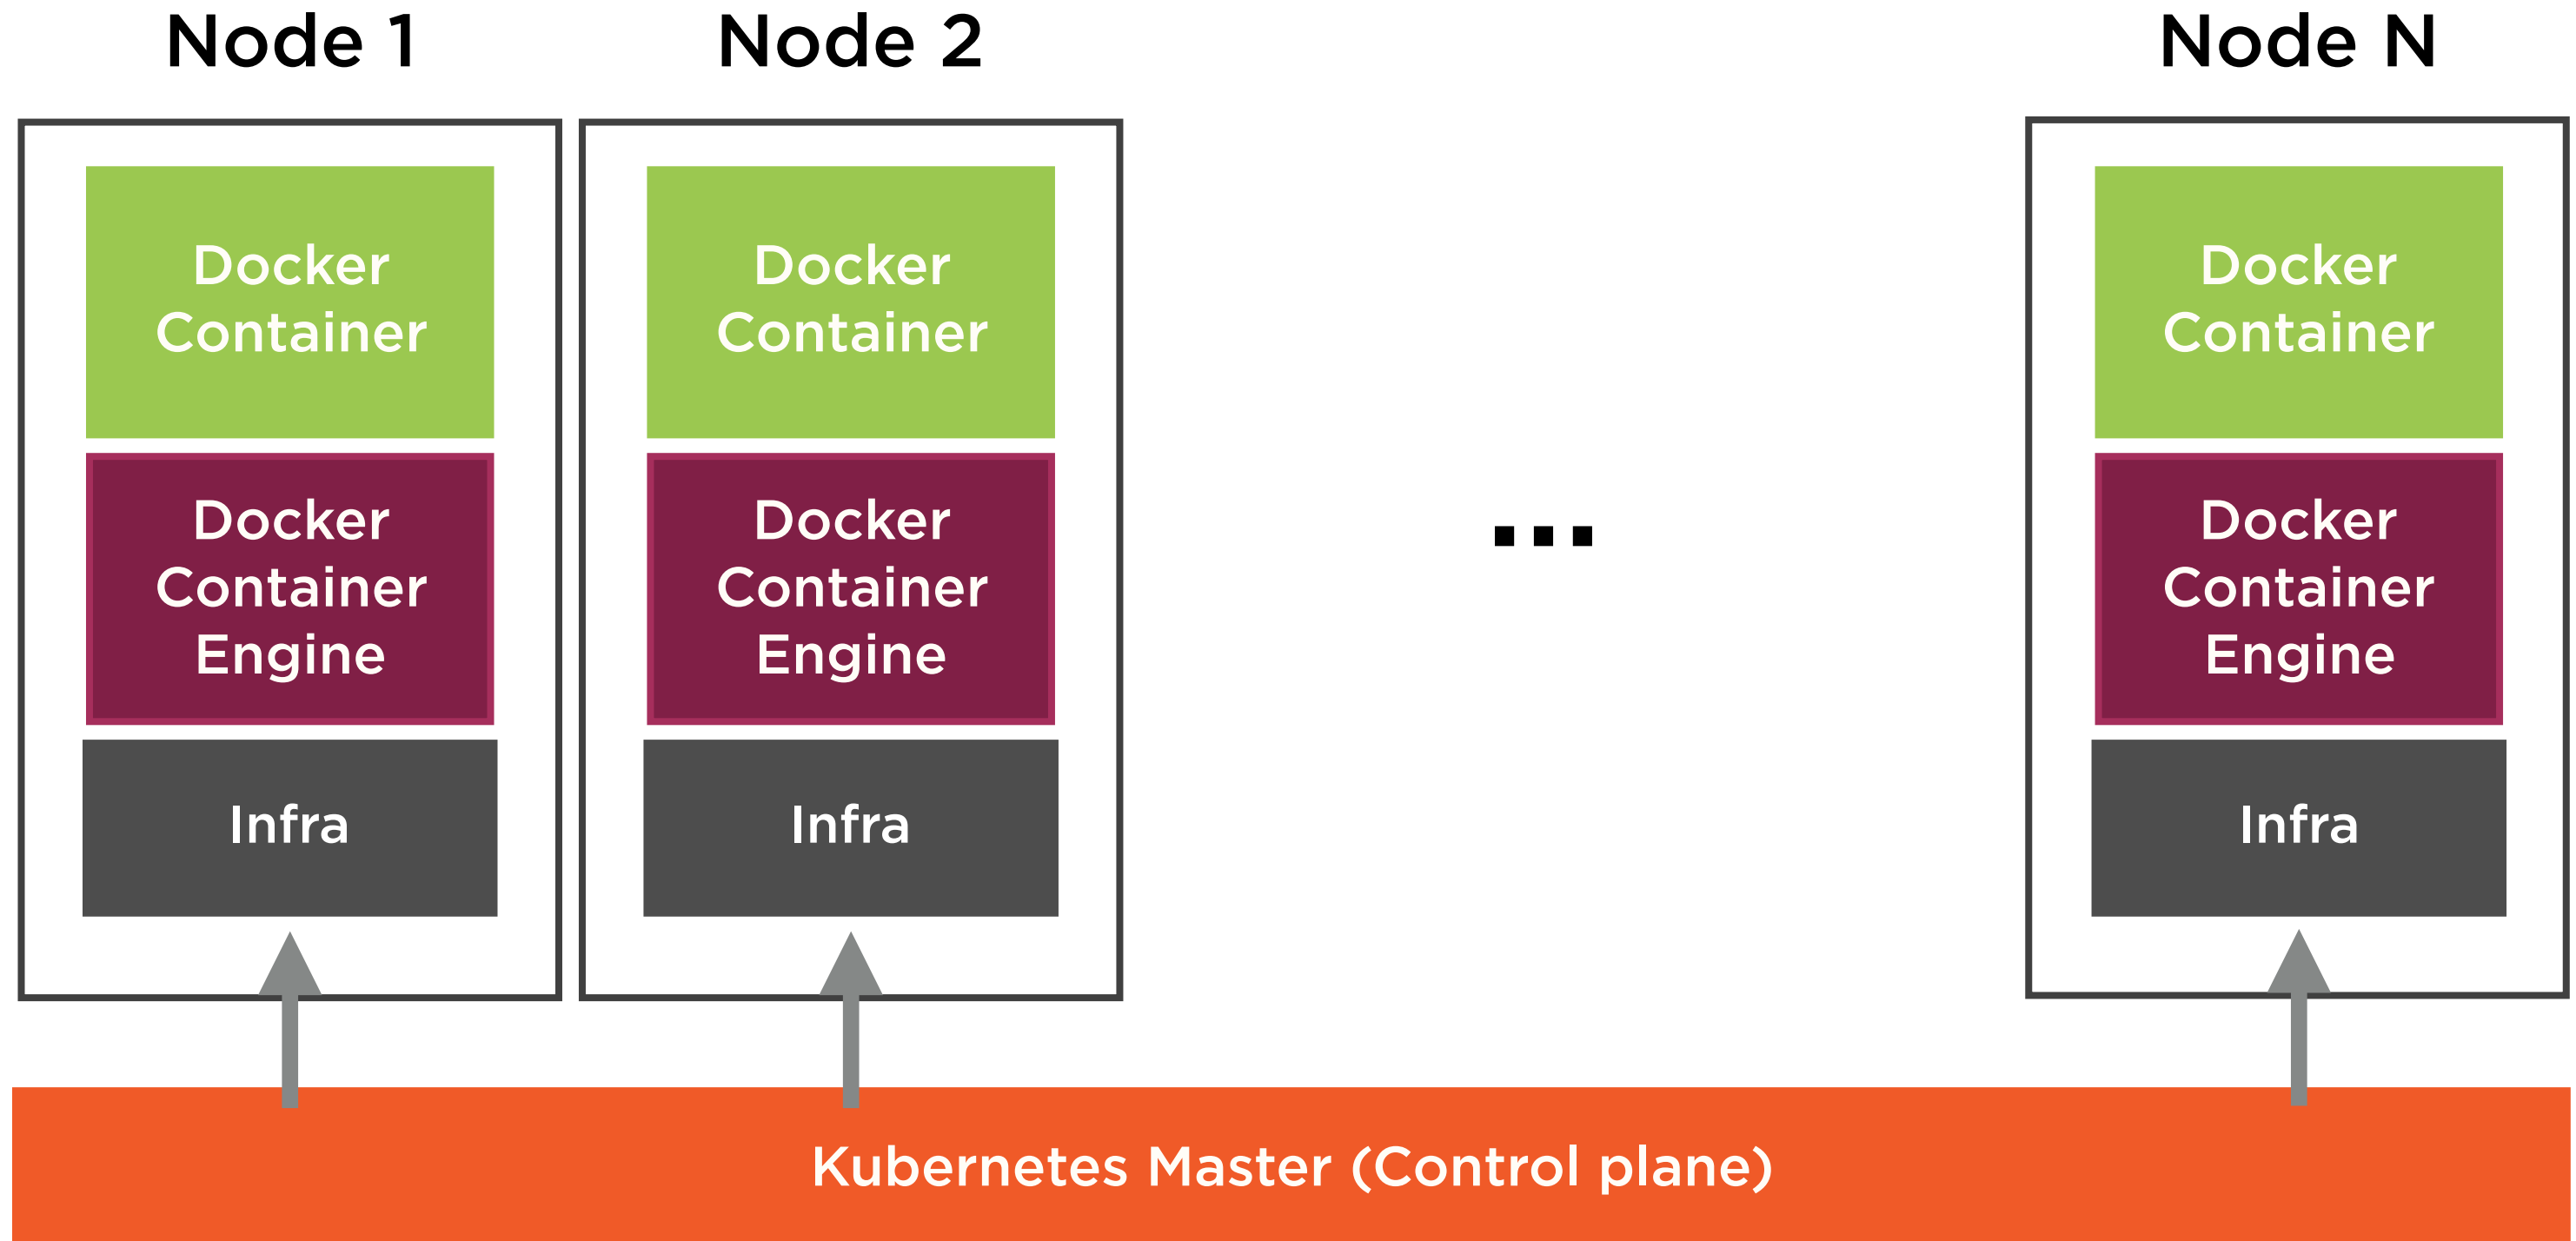
**Potentially thousands of containers on hundreds of VMs**



# Isolated Container Deployments



# Kubernetes: Cluster Orchestration



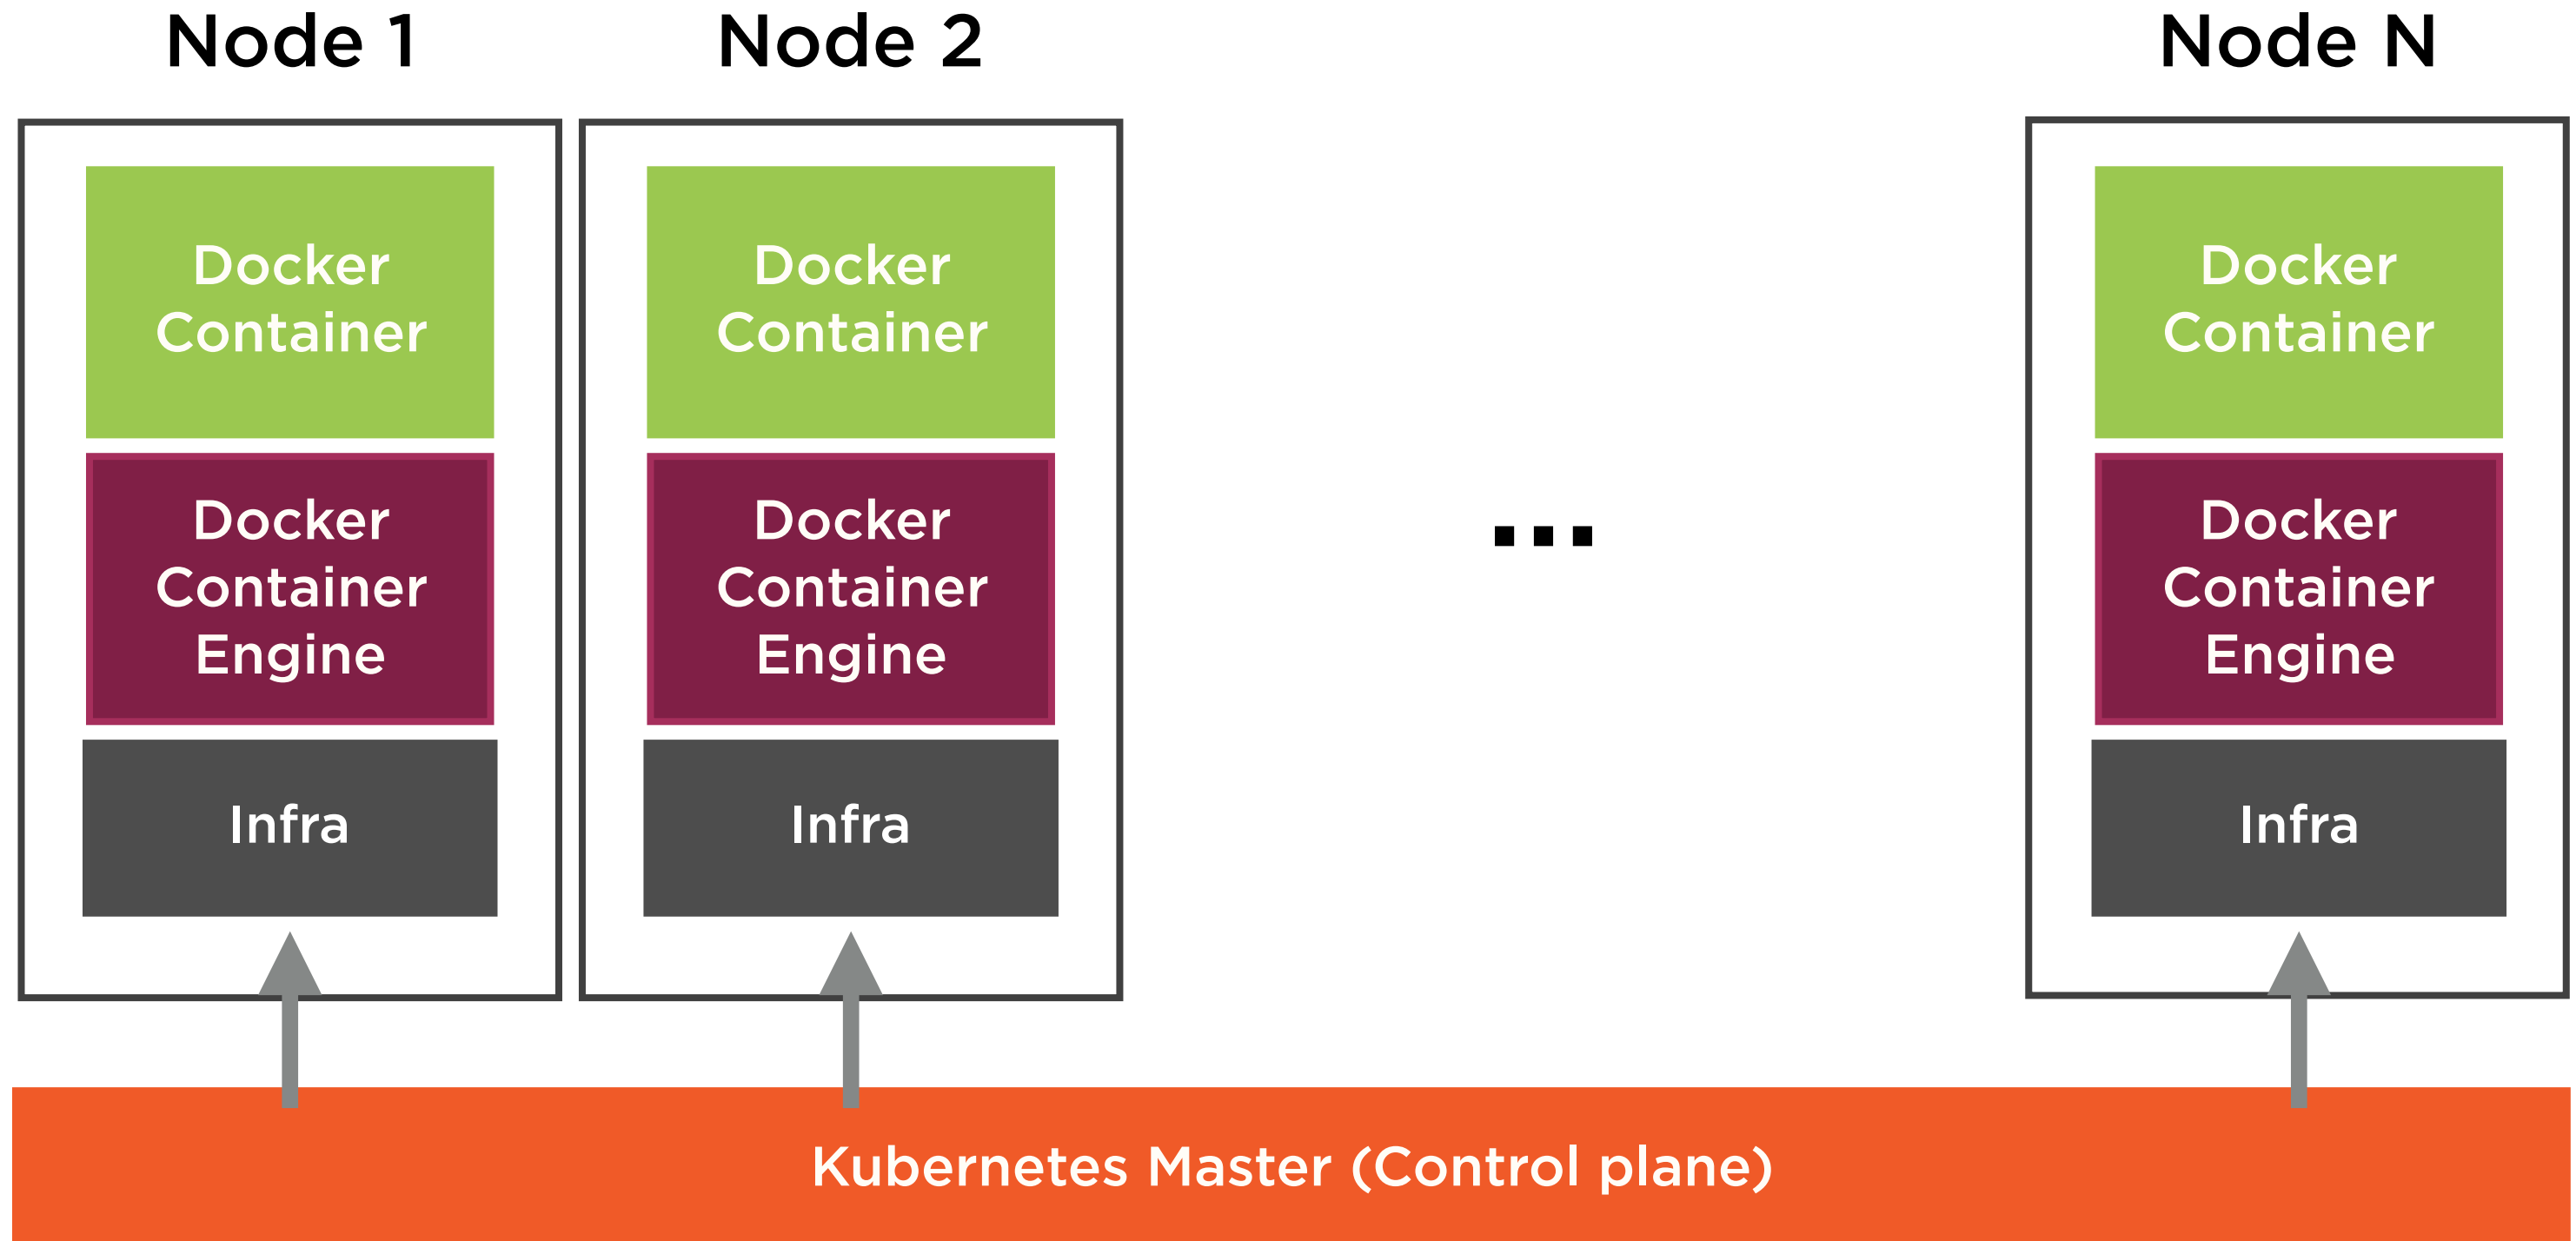


Kubernetes does not interact directly  
with containers


Instead it uses a number of higher-level  
entities referred to as **objects**



# Kubernetes: Cluster Orchestration



# Kubernetes: Cluster Orchestration




# Kubernetes: Containers Run Within Pods

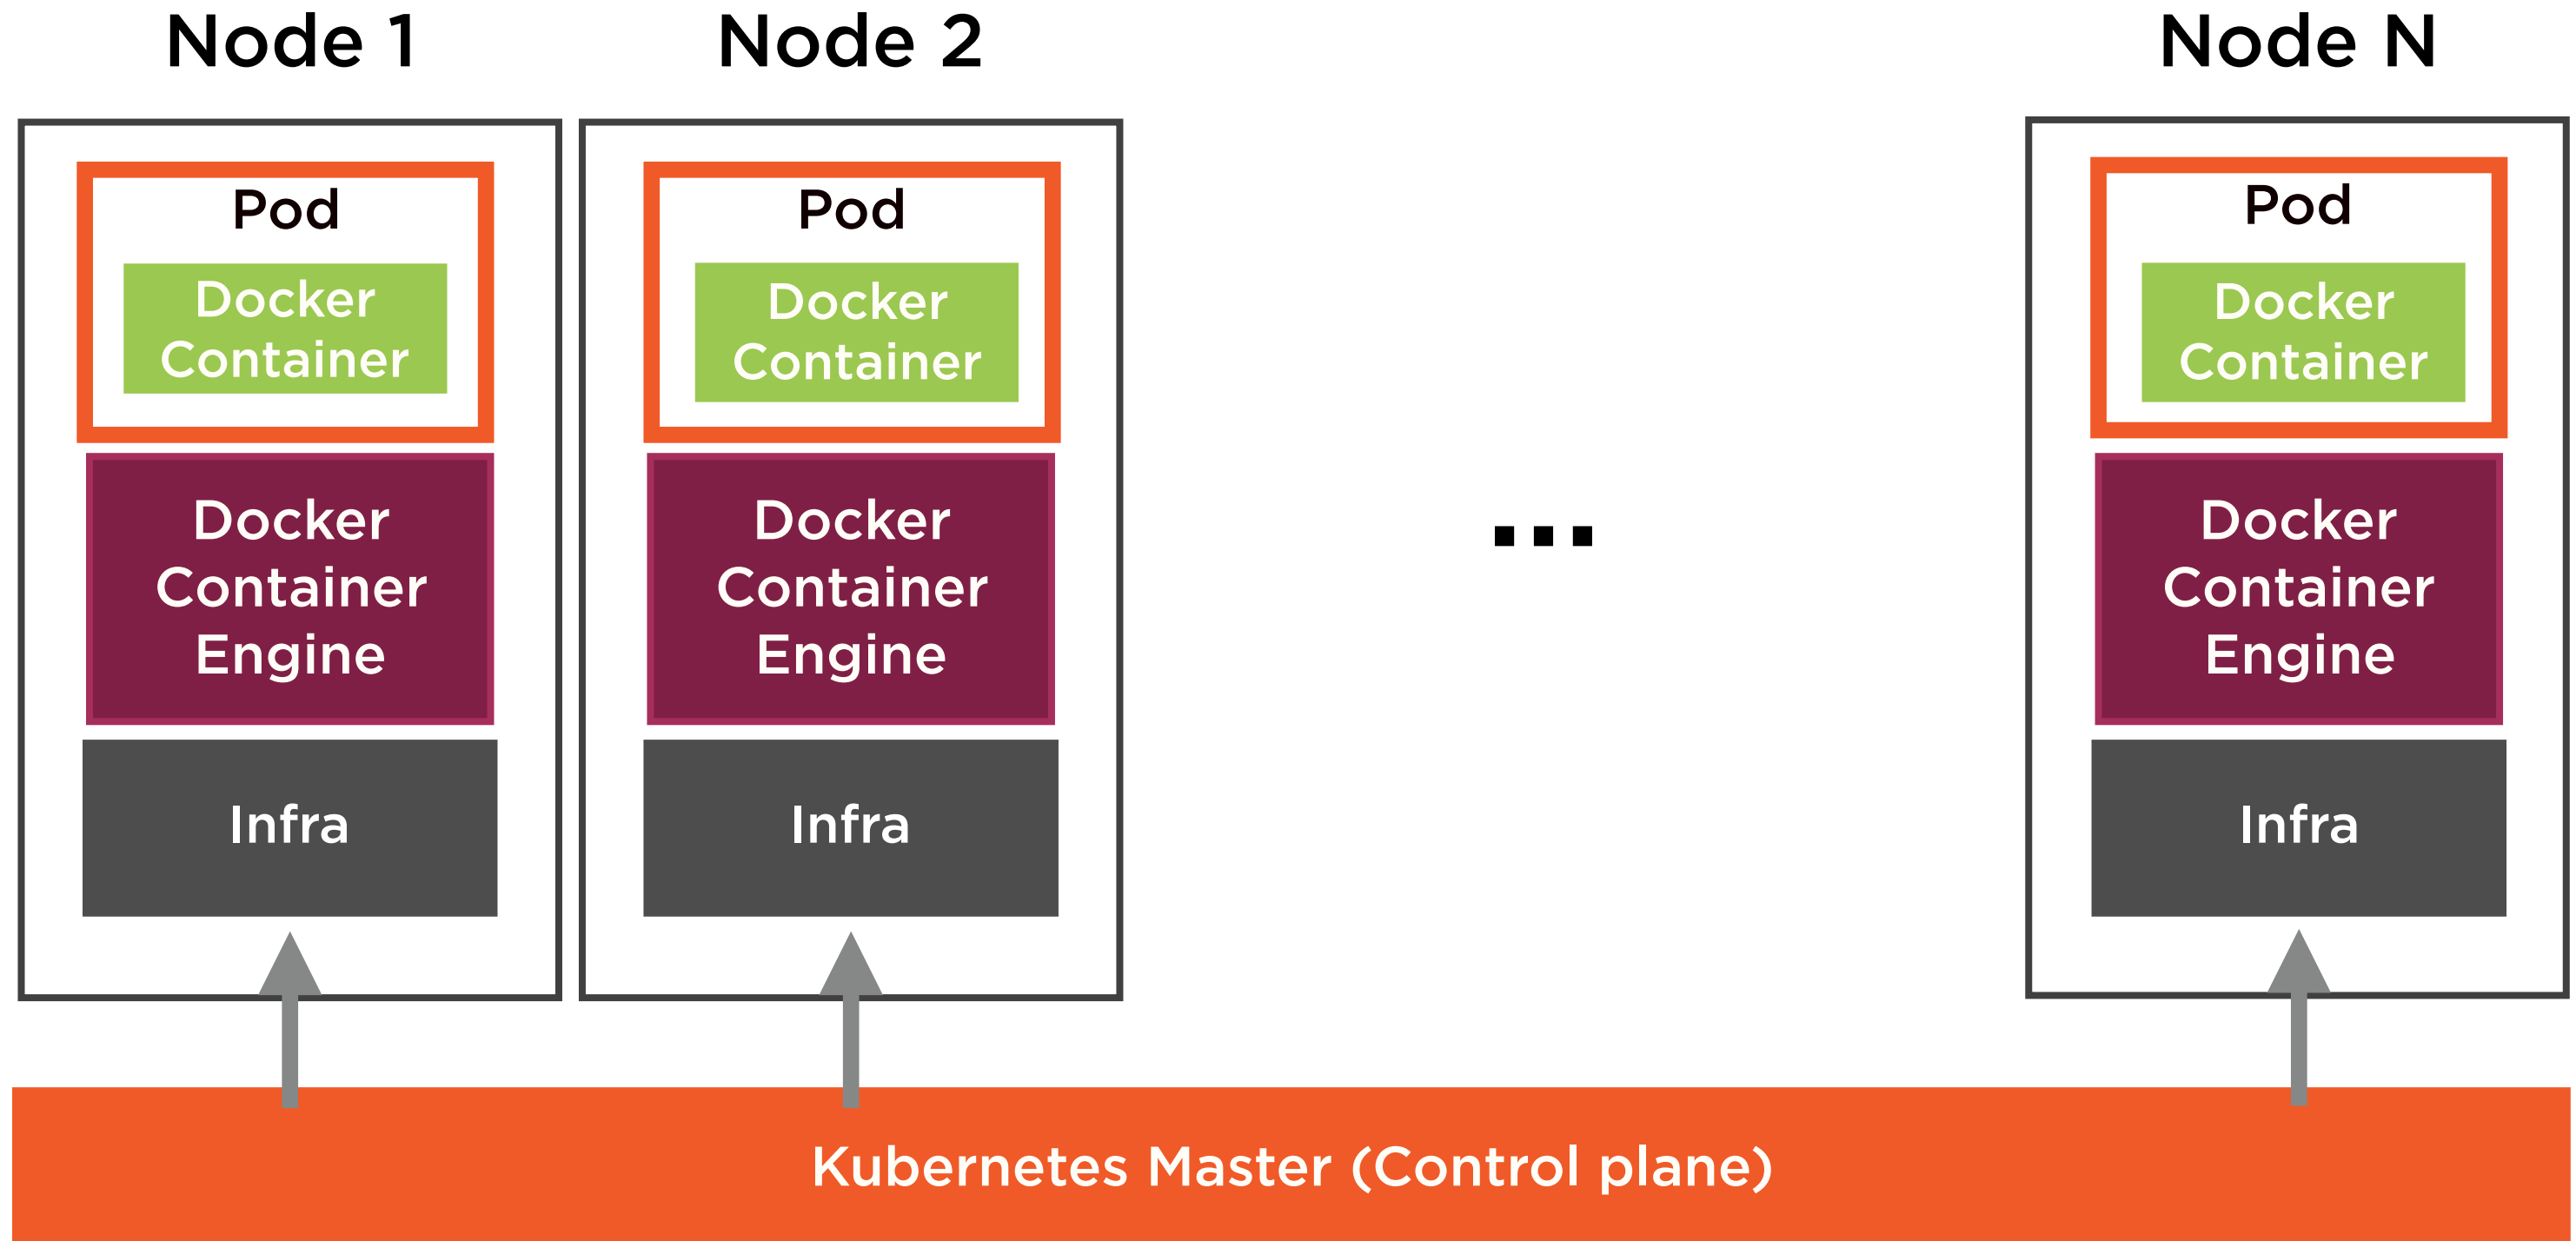
**Node 1**

**Node 2**

**Node N**



# Kubernetes: Cluster Orchestration



# Pods as Atomic Units

## Container deployment

All containers in pod are deployed,  
or none are

## Node association

Entire pod is hosted on the same  
node

**Pod is atomic unit of deployment in Kubernetes**



# Pods on Kubernetes Nodes



**Encapsulates one or more containers**

**Pods run on nodes**

**Nodes are controlled by master**

**In GKE**

- Nodes are GCE VM instances
- Master is managed by GKE service

# Pods on Kubernetes Nodes



**Can not run a container without enclosing pod**

**Pods provide isolation between containers**

**Pod acts as sandbox for enclosed containers**

**Multi-container pods are possible**

- tightly-coupled
- not usually recommended






The **Pod object** is the lowest level of abstraction around a container

Every object is associated with a specification file which represents the **desired end state** of the object



# YAML Specification Files



**Current State**


The current state of the object

**Desired State**

The end state of the object



# YAML Specification Files



Controllers in the Kubernetes cluster run reconciliation loops to get the actual state to match the desired state



# Pod Specification File

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    test: liveness
    name: liveness-http
spec:
  containers:
  - args:
    - /server
    image: k8s.gcr.io/liveness
    livenessProbe:
      httpGet:
        # when "host" is not defined, "PodIP" will be used
        # host: my-host
        # when "scheme" is not defined, "HTTP" scheme will be used. Only "HTTP" and "HTTPS" are allowed
        # scheme: HTTPS
        path: /healthz
        port: 8080
        httpHeaders:
        - name: X-Custom-Header
          value: Awesome
        initialDelaySeconds: 15
        timeoutSeconds: 1
    name: liveness
```

# Pod Specification File

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    test: liveness
  name: liveness-http
spec:
  containers:
  - args:
    - /server
    image: k8s.gcr.io/liveness
    livenessProbe:
      httpGet:
        # when "host" is not defined, "PodIP" will be used
        # host: my-host
        # when "scheme" is not defined, "HTTP" scheme will be used. Only "HTTP" and "HTTPS" are allowed
        # scheme: HTTPS
        path: /healthz
        port: 8080
        httpHeaders:
        - name: X-Custom-Header
          value: Awesome
        initialDelaySeconds: 15
        timeoutSeconds: 1
      name: liveness
```

Which container image(s)?

Available on which port?

Using the Google Kubernetes Engine almost completely eliminates the need to explicitly configure YAML files

**Simply use the web console or the gcloud command line utility**





# Pod Spec

**Pods are automatically assigned unique IP addresses**

**For each container**

- Container image with source repository
- Port on which container will be available

**Pods can specify shared storage volumes**

**Pod can contain multiple containers**





# Multi-container Pods

**Avoid unless necessary**

**Tight coupling, no isolation**

**No independent scaling possible**

**Prefer service-oriented architecture to monolithic design**





# Pod Status



## Pending:

- Request accepted, but not yet fully created

## Running:

- Pod bound to node, all containers started

## Succeeded:

- All containers terminated successfully (will not be restarted)

## Failed:

- All containers have terminated, and at least one failed

## Unknown:

- Pod status could not be queried



# Kubernetes for Container Orchestration

---

# Limitation of Standalone Pods



## No autohealing

- Crashed pods won't restart automatically
- Need higher level orchestration

## No scaling or autoscaling

- Overloaded pods don't spawn more automatically
- Need higher level orchestration

## No load balancing

- Pod IP addresses are ephemeral
- Pods can't share load automatically
- Need higher level orchestration



# Higher-level Abstractions



## **ReplicaSet, ReplicationController**

- Scaling and healing

## **Deployment**

- Versioning and rollback

## **Service**

- Static (non-ephemeral) IP addresses
- Stable networking

## **Persistent volumes**

- Non-ephemeral storage



# Kubernetes as Orchestrator



## **Fault tolerance**

- Recover from pod or node failures

## **Autohealing**

- Crashed containers restart

## **Scaling**

- ReplicaSets for multiple copies of the pod



# Kubernetes as Orchestrator



## Autoscaling

- Scale up and scale down based on load
- Horizontal pod autoscalers

## Isolation

- Containers run inside pods

## Load balancing

- Distribute traffic to containers



# Kubernetes as Orchestrator



**Master**

**Worker nodes**

**Job scheduling**

**Resource allocation**

**Comparing actual and desired state**



# Kubernetes as Orchestrator



**Docker container engine : Java Runtime**

**Docker containers : Jar files**

**Kubernetes : Hadoop**





# Kubernetes as Orchestrator

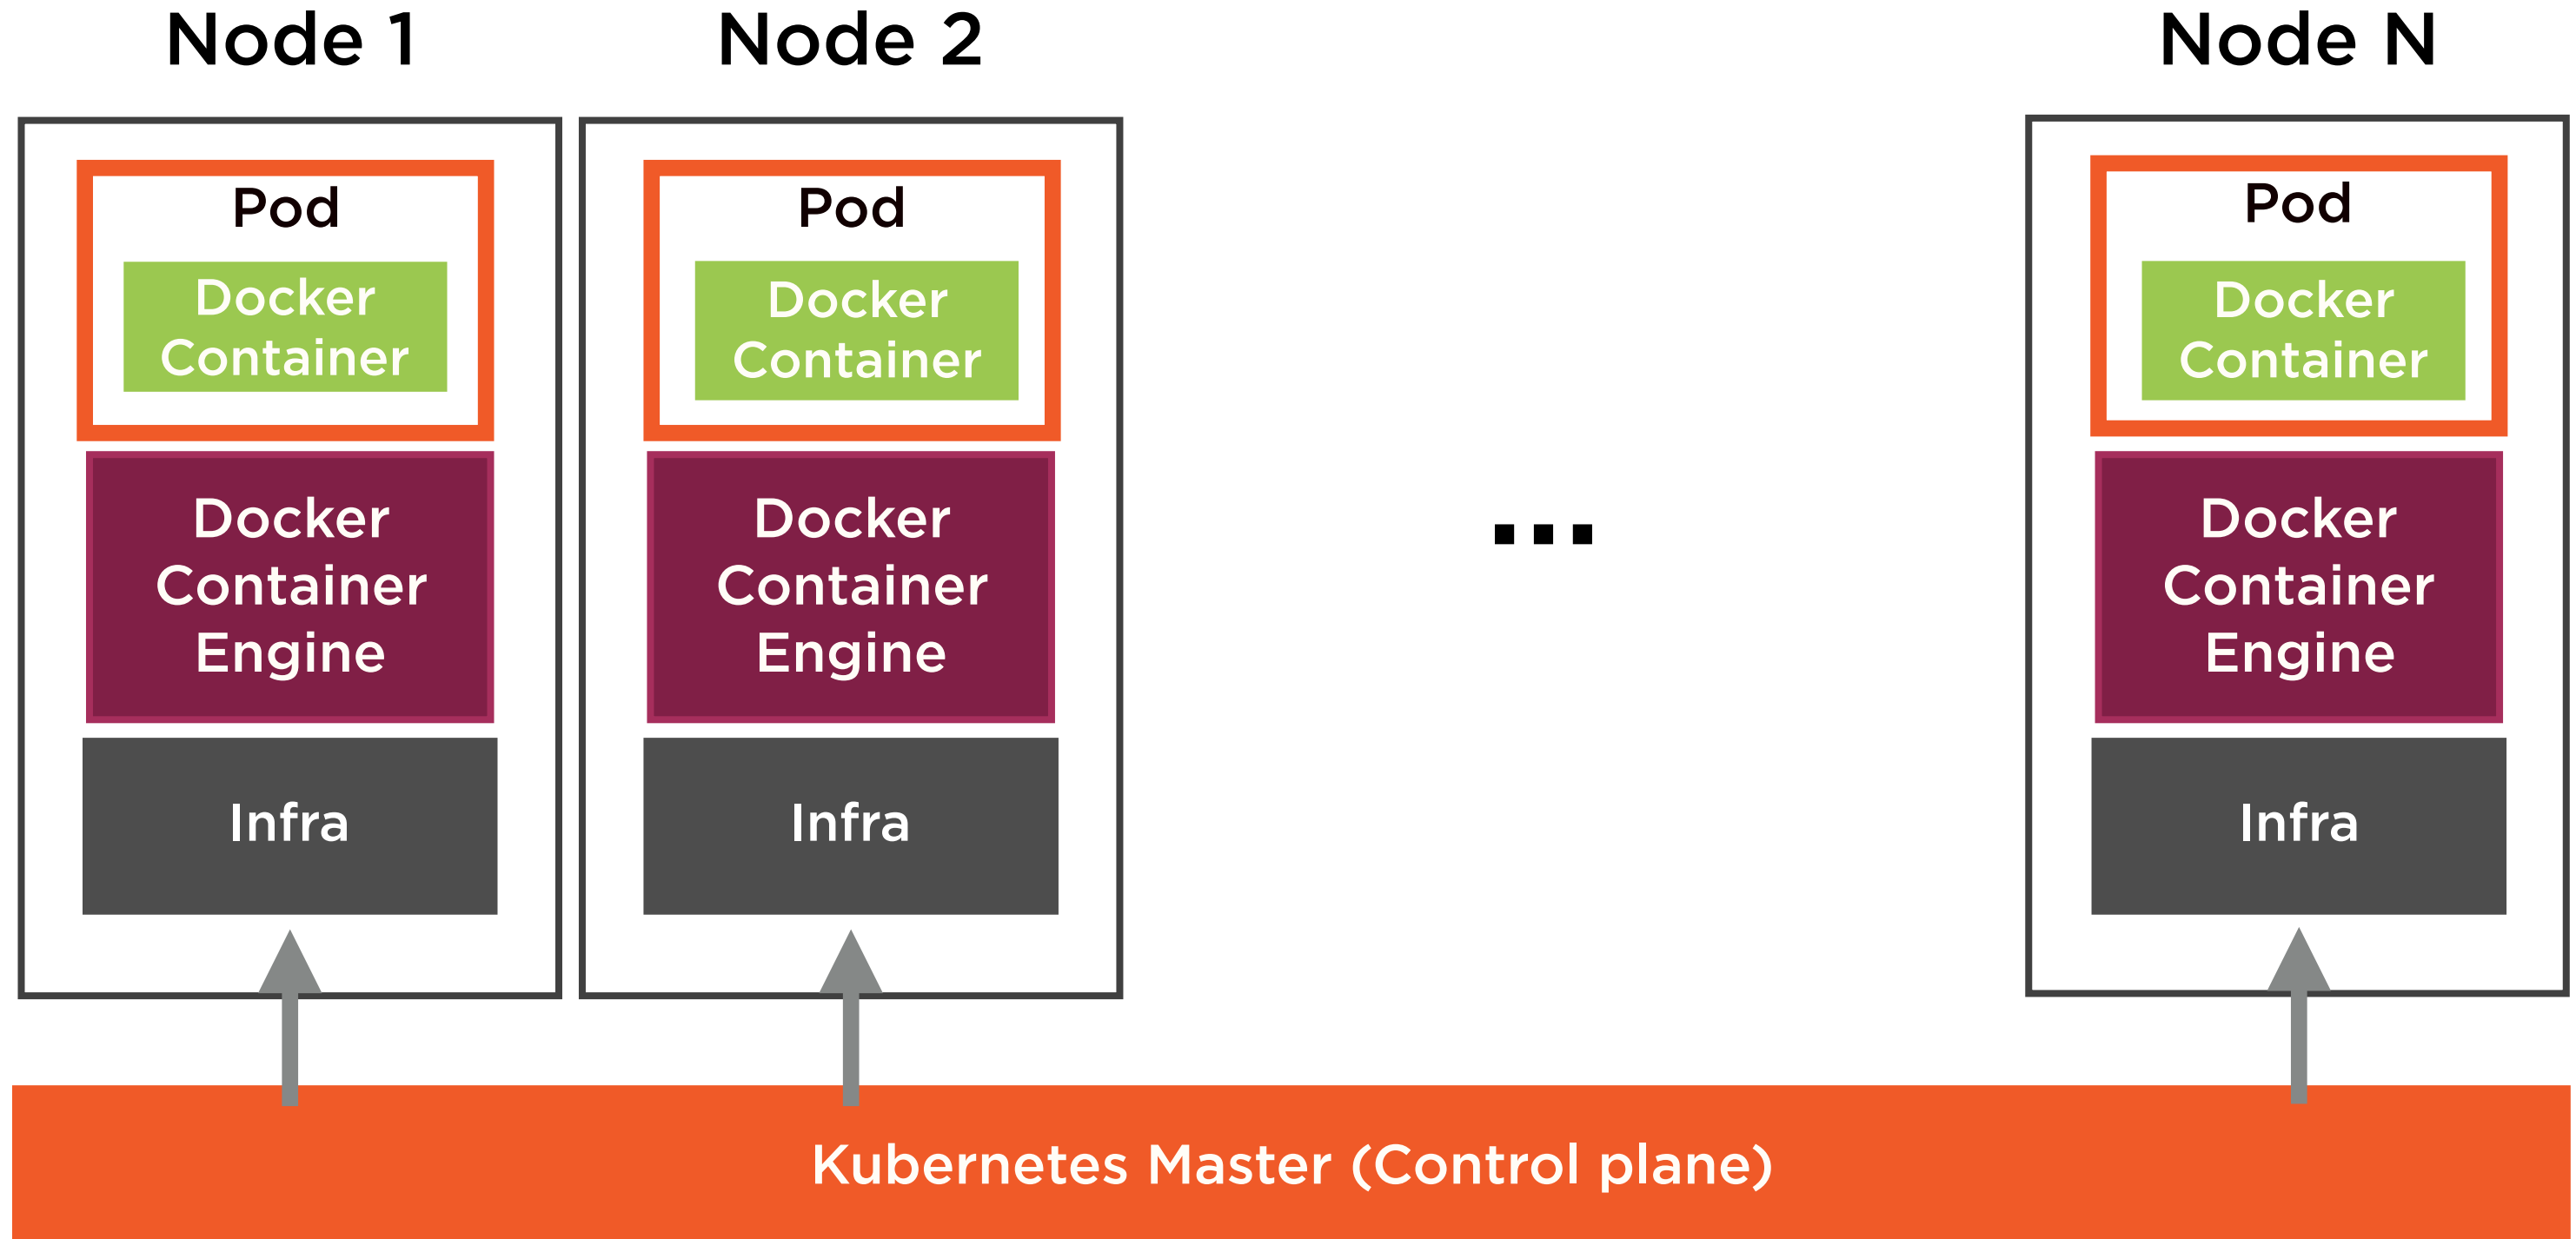


## **User interacts with Kubernetes**

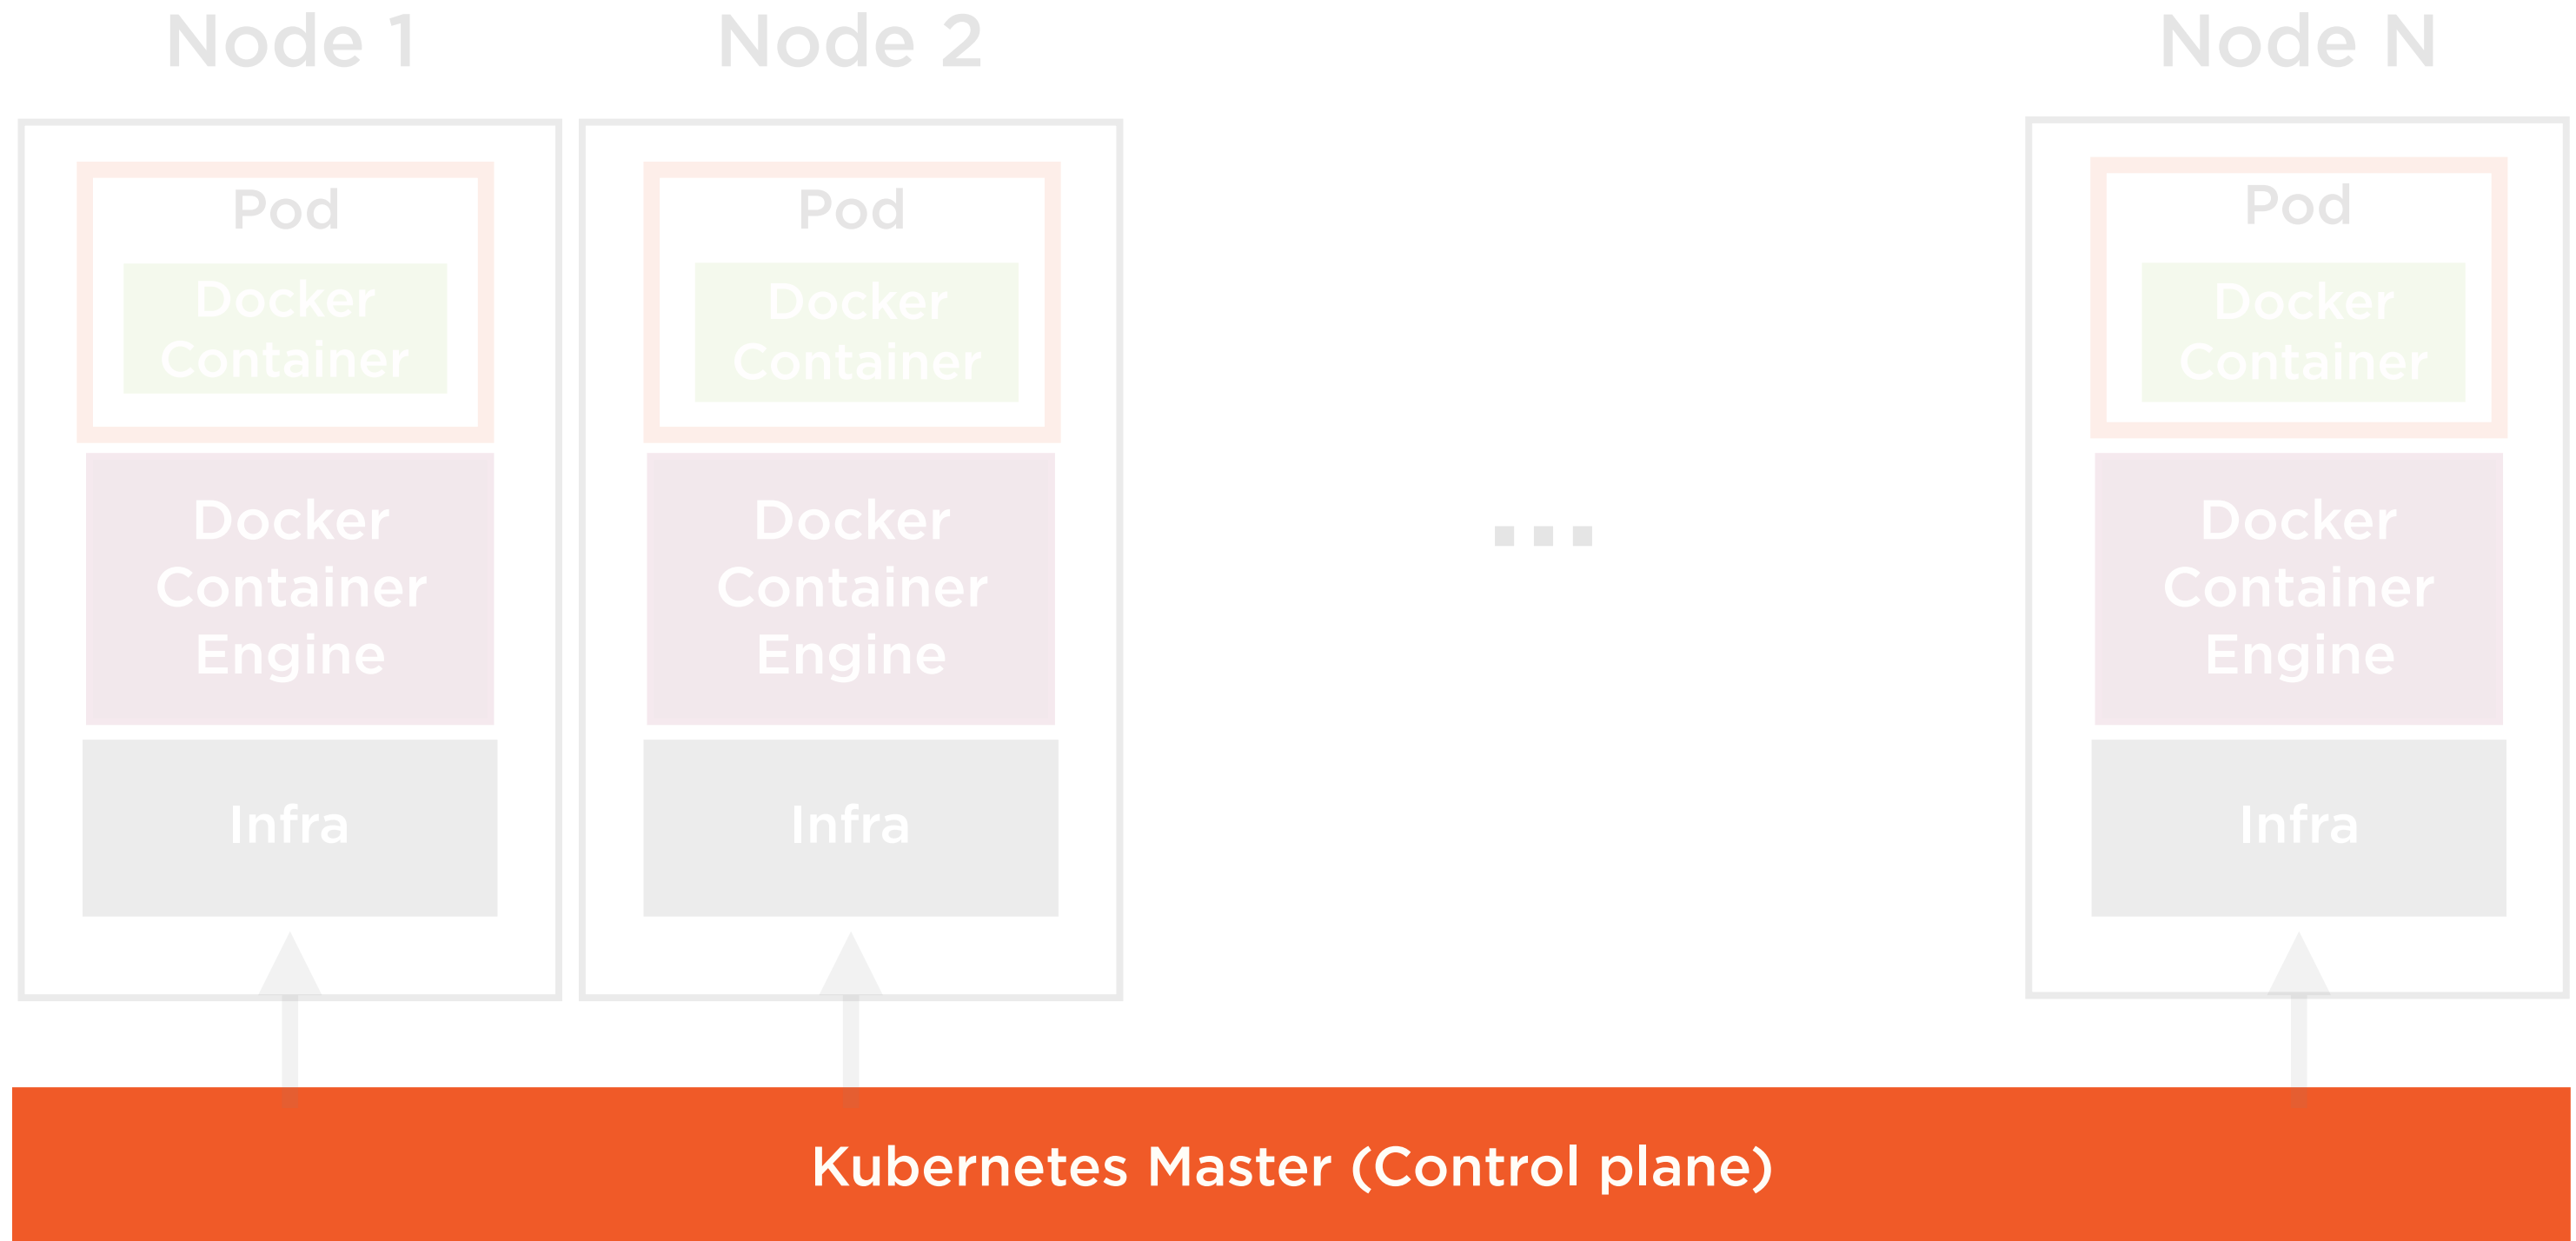
- kubectl and other command line tools
- YAML config files

## **Kubernetes relays and orchestrates pods**

# Kubernetes: Cluster Orchestration



# Kubernetes: Cluster Orchestration



# Replication and Deployment

---



# Higher-level Abstractions



## **ReplicaSet, ReplicationController**

- Scaling and healing

## **Deployment**

- Versioning and rollback

## **Service**

- Static (non-ephemeral) IP addresses
- Stable networking

## **Persistent volumes**

- Non-ephemeral storage



# Higher-level Abstractions



## **ReplicaSet, ReplicationController**

- Scaling and healing

## Deployment

- Versioning and rollback

## Service


- Static (non-ephemeral) IP addresses
- Stable networking

## Persistent volumes

- Non-ephemeral storage



# Kubernetes: Cluster Orchestration



# ReplicaSet



Multiple copies of the pod are managed together using a ReplicaSet

Self-healing and autoscaling for our pods





Multiple instances of the pod are created and deployed to clusters

Replicas are represented by a  
ReplicaSet object





# ReplicaSet

**If pod crashes, ReplicaSet will start a new one**

**Key to scaling and healing**

**All pods are replicas of each other**





# ReplicaSet

**Loosely coupled with pods**

**A ReplicaSet object will govern all pods that match its label selector**

**Users must ensure no conflicts, orphans**



# The ReplicaSet Object

**Multiple identical pods which  
are replicas of each other**



# Higher-level Abstractions



## ReplicaSet, ReplicationController

- Scaling and healing

## Deployment

- Versioning and rollback

## Service


- Static (non-ephemeral) IP addresses
- Stable networking

## Persistent volumes

- Non-ephemeral storage



# Deployment



Works with ReplicaSets to add versioning, rollback and other advanced deployment functionality

# The Deployment Object

**Adds on deployment and  
rollback functionality**





# Deployment Objects

**Easy to push out new version of container**

**Triggers creation of new ReplicaSet and new containers**

**Pods in old ReplicaSet gradually reduced to zero**







# Deployment Rollbacks

**Every change to a Deployment object triggers creation of a new revision**

**Trivial to rollback to previous revision**

**Offers versioning support**





# Use Cases of Deployments

**Manual scaling: Edit number of replicas**

**Autoscaling: Use HPA with deployment as target**

**Can pause/resume deployments midway**

**Can monitor deployment status**



# Services

---



# Limitation of Standalone Pods



## No autohealing

- Crashed pods won't restart automatically
- Need higher level orchestration

## No scaling or autoscaling

- Overloaded pods don't spawn more automatically
- Need higher level orchestration

## No load balancing

- Pod IP addresses are ephemeral
- Pods can't share load automatically
- Need higher level orchestration



# Higher-level Abstractions



## **ReplicaSet, ReplicationController**

- Scaling and healing

## **Deployment**

- Versioning and rollback

## **Service**

- Static (non-ephemeral) IP addresses
- Stable networking

## **Persistent volumes**

- Non-ephemeral storage





# Higher-level Abstractions

## ReplicaSet, ReplicationController

- Scaling and healing

## Deployment

- Versioning and rollback

## Service

- Static (non-ephemeral) IP addresses
- Stable networking

## Persistent volumes

- Non-ephemeral storage



# Ephemeral IP Addresses

**Containers expose ports in pod spec**

**Pod IP addresses are ephemeral**

**Poses a problem for clients**

**How are they to know where to send  
request?**



Services provide stable IP  
addresses for external  
connections and load balancing





# Service Objects

**Provides stable (non-ephemeral) IP address**

**Connects to set of back-end pods**

**Set of pods changes dynamically**

- Logically selected via label selector

**Front-end IP remains unchanged**

**Basic load balancing too**





# Endpoint Object

**Each service has associated endpoint object**

**Dynamic list of pods selected by service**

**Populated by Kubernetes**

**Dynamically updates as pods are created/deleted**




# Volume Abstractions

---



# Storage with Containers



## **On disk files within a container**

- Only accessible to the container itself
- Ephemeral: is lost when the container stops or crashes

## **Volume abstractions**

- A directory accessible to all containers in a pod
- Specified in the PodSpec
- Have the same lifetime as the enclosing pod

For durable storage use  
persistent volumes

The volume is preserved even  
when the pod is removed and  
can be handed off to another pod



# Higher-level Abstractions



## **ReplicaSet, ReplicationController**

- Scaling and healing

## **Deployment**

- Versioning and rollback

## **Service**

- Static (non-ephemeral) IP addresses
- Stable networking

## **Persistent volumes**

- Non-ephemeral storage





# Higher-level Abstractions

## ReplicaSet, ReplicationController

- Scaling and healing

## Deployment

- Versioning and rollback

## Service

- Static (non-ephemeral) IP addresses
- Stable networking

## Persistent volumes

- Non-ephemeral storage

# Volumes

## Permanence

Storage that outlives an individual container or individual pod


## Shared state

Safe mechanism for containers in pod to share state

**Persistent volumes** are volumes whose life is not tied to an individual pod







# Types of Volumes

**configMap**


**emptyDir**

**gitRepo**

**secret**

**hostPath**





# Using Volumes

**Define volume in pod spec**

**Have each container mount volume**

**Each container mounts independently**

**At different path**



# Volumes and Persistent Volumes

## Volumes


**Storage abstraction with life longer than individual container inside pod**

**Life tied to life of pod**

## Persistent Volumes

**Storage abstraction with life longer than individual pod**

**Life not tied to life of pod**



# Cloud-specific Persistent Volumes


**awsElasticBlockStore**

**azureDisk**

**azureFile**

**gcePersistentDisk**





# Persistent Volumes

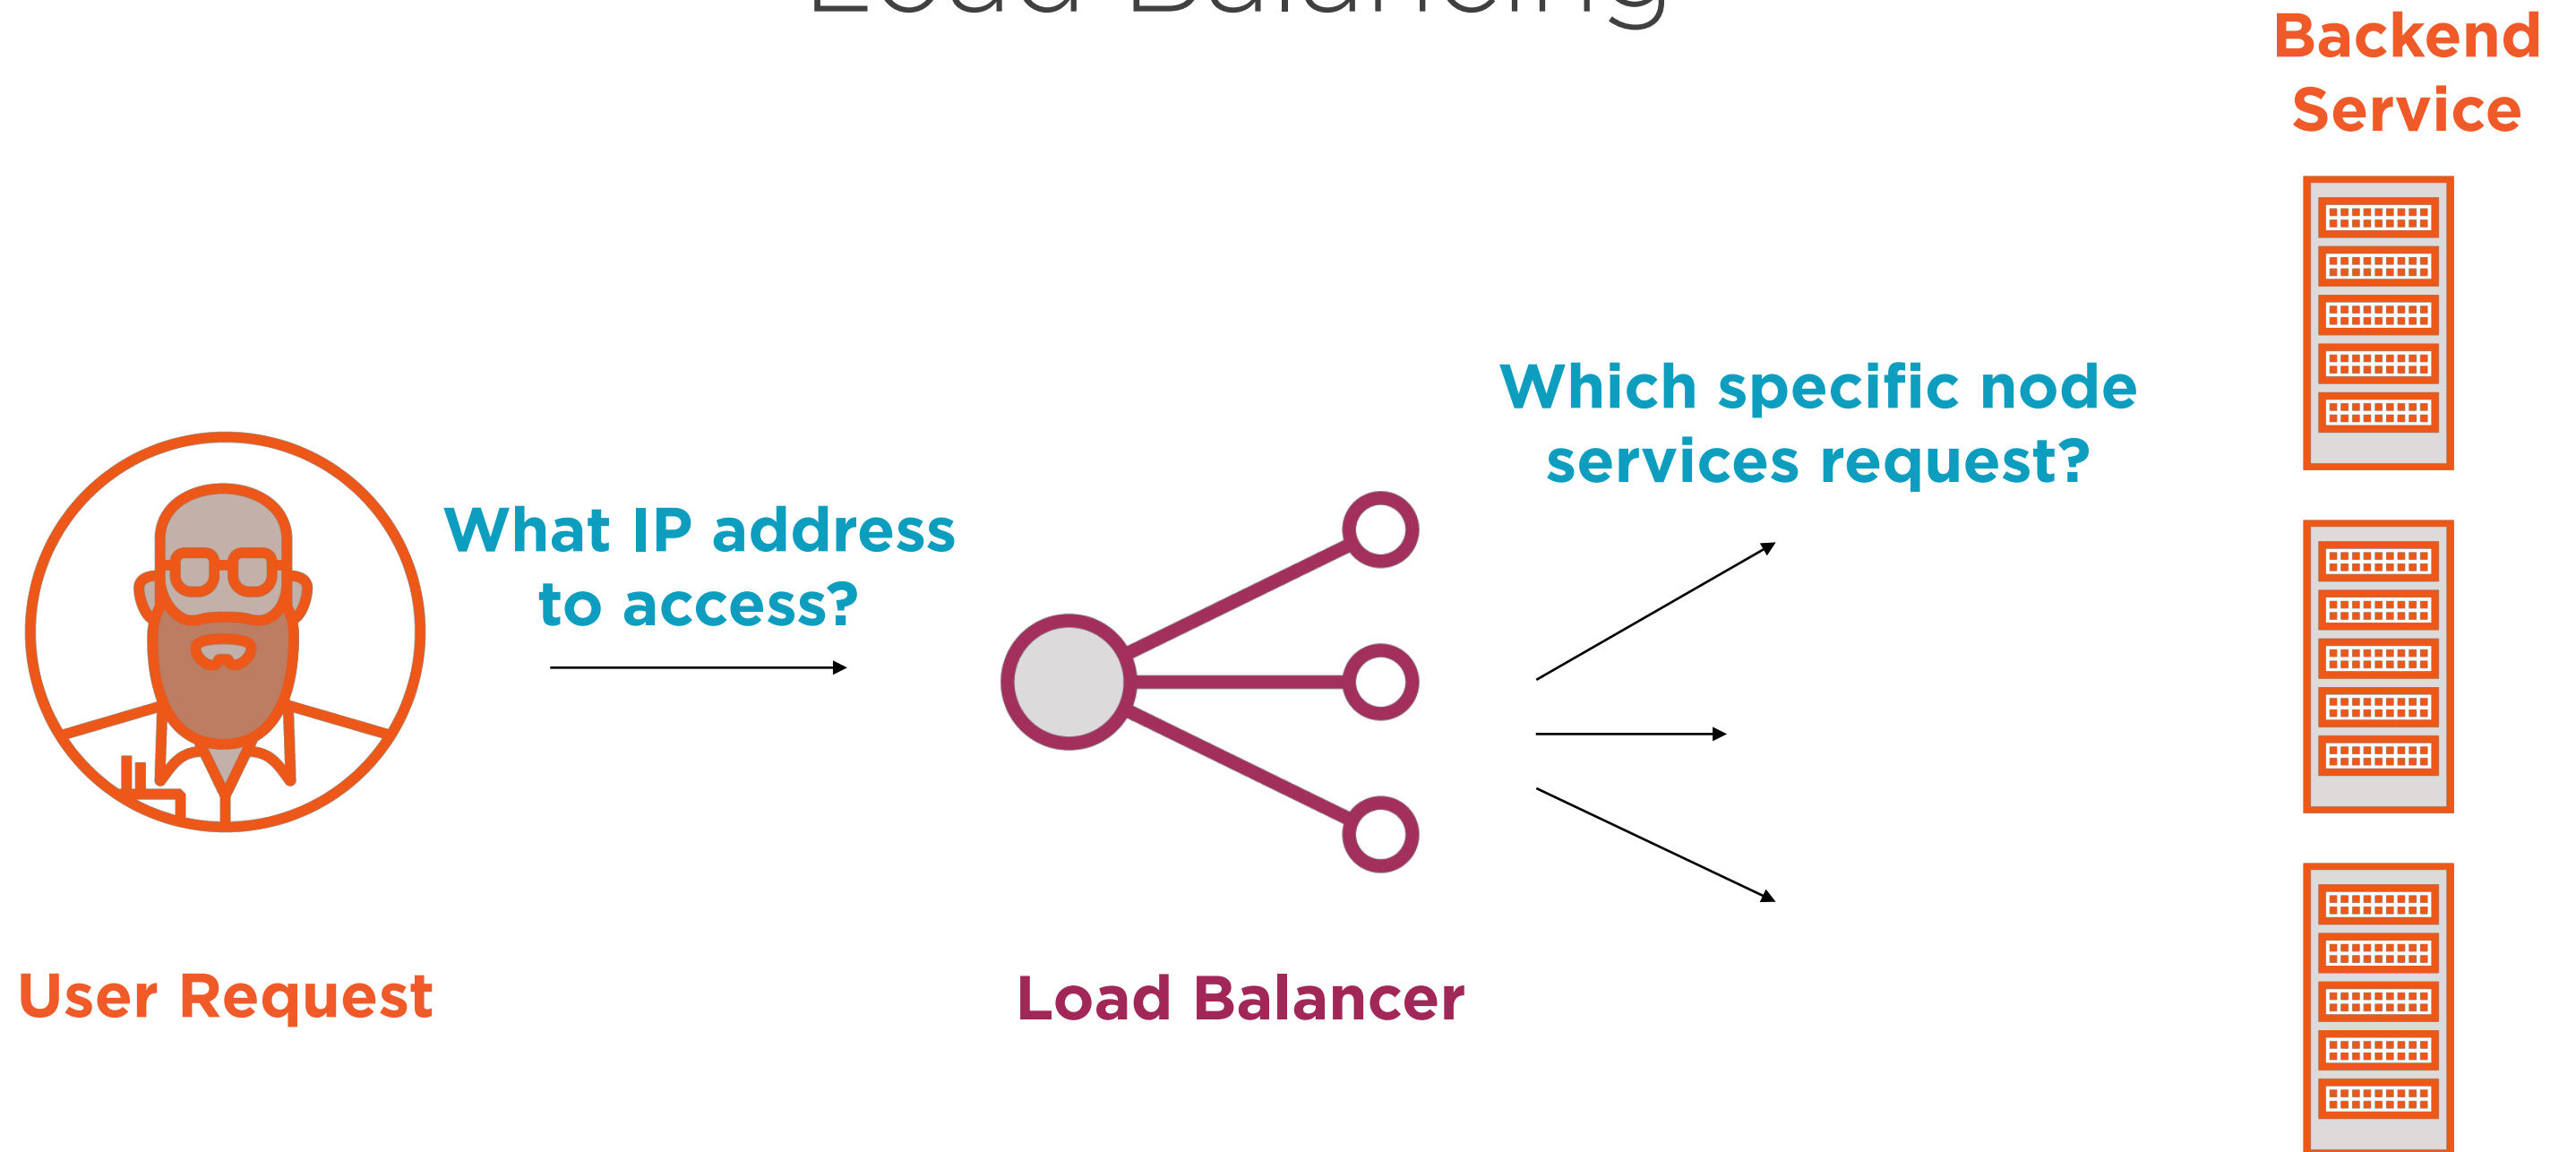
## Two types of provisioning

- Static: Administrator pre-creates volume
- Dynamic: Containers gain access by filing PersistentVolumeClaim

# Load Balancing


---

# Load Balancing



Load Balancer to the rescue





# Load Balancers

**Stable front-end IP**

**Forwarding rules to funnel traffic**

**Connect to backend service**

**Distribute load intelligently**

**Health checks to avoid unhealthy instances**





Exposing a service on the GKE  
allows the option of configuring a  
load balancer to distribute traffic



# Ingress

---


# Ingress Object

Kubernetes object defining a collection of rules that allow inbound connections to reach cluster services. On GKE, a single ingress object can control access to multiple services




**A single service can  
expose an IP address for  
access**

Ingress



With multiple services it  
makes sense to have  
rules defined using an  
ingress object





# Ingress Objects

**Can be configured to support**

- Externally-reachable URLs
- Load balancing
- SSL termination
- Name-based virtual hosting





# Ingress Objects

**GKE fully supports both ingress and load balancer objects**

**On other platforms need to do some additional work**

**(Need to configure ingress object as pod)**





# GKE Ingress Objects

**GKE clusters have HttpLoadBalancing add-on enabled**

**Causes additional controller to run on master**

**This controller supports HTTP(S) load balancing for ingress objects**





# StatefulSets and DaemonSets

---


# StatefulSets

**A set of pods, similar to ReplicaSet**

**Important difference from ReplicaSet**

- Pods created unique
- Identified by name
- Not interchangeable
- Always associated with persistent volume





# DaemonSet


**Manages groups of replicated pods**

**Attempts to keep one pod per node**

**Across all nodes or a subset**

**As nodes added, pods created too**

**As nodes removed, pods are garbage collected**



# DaemonSet

**Cluster storage daemons**

**Log collection daemons**

**Node monitoring daemons**




# Horizontal Pod Autoscaler

---



# Autoscaling with Kubernetes



## **Autoscaling**

Compute capacity automatically changes with changing need




## **Autohealing**

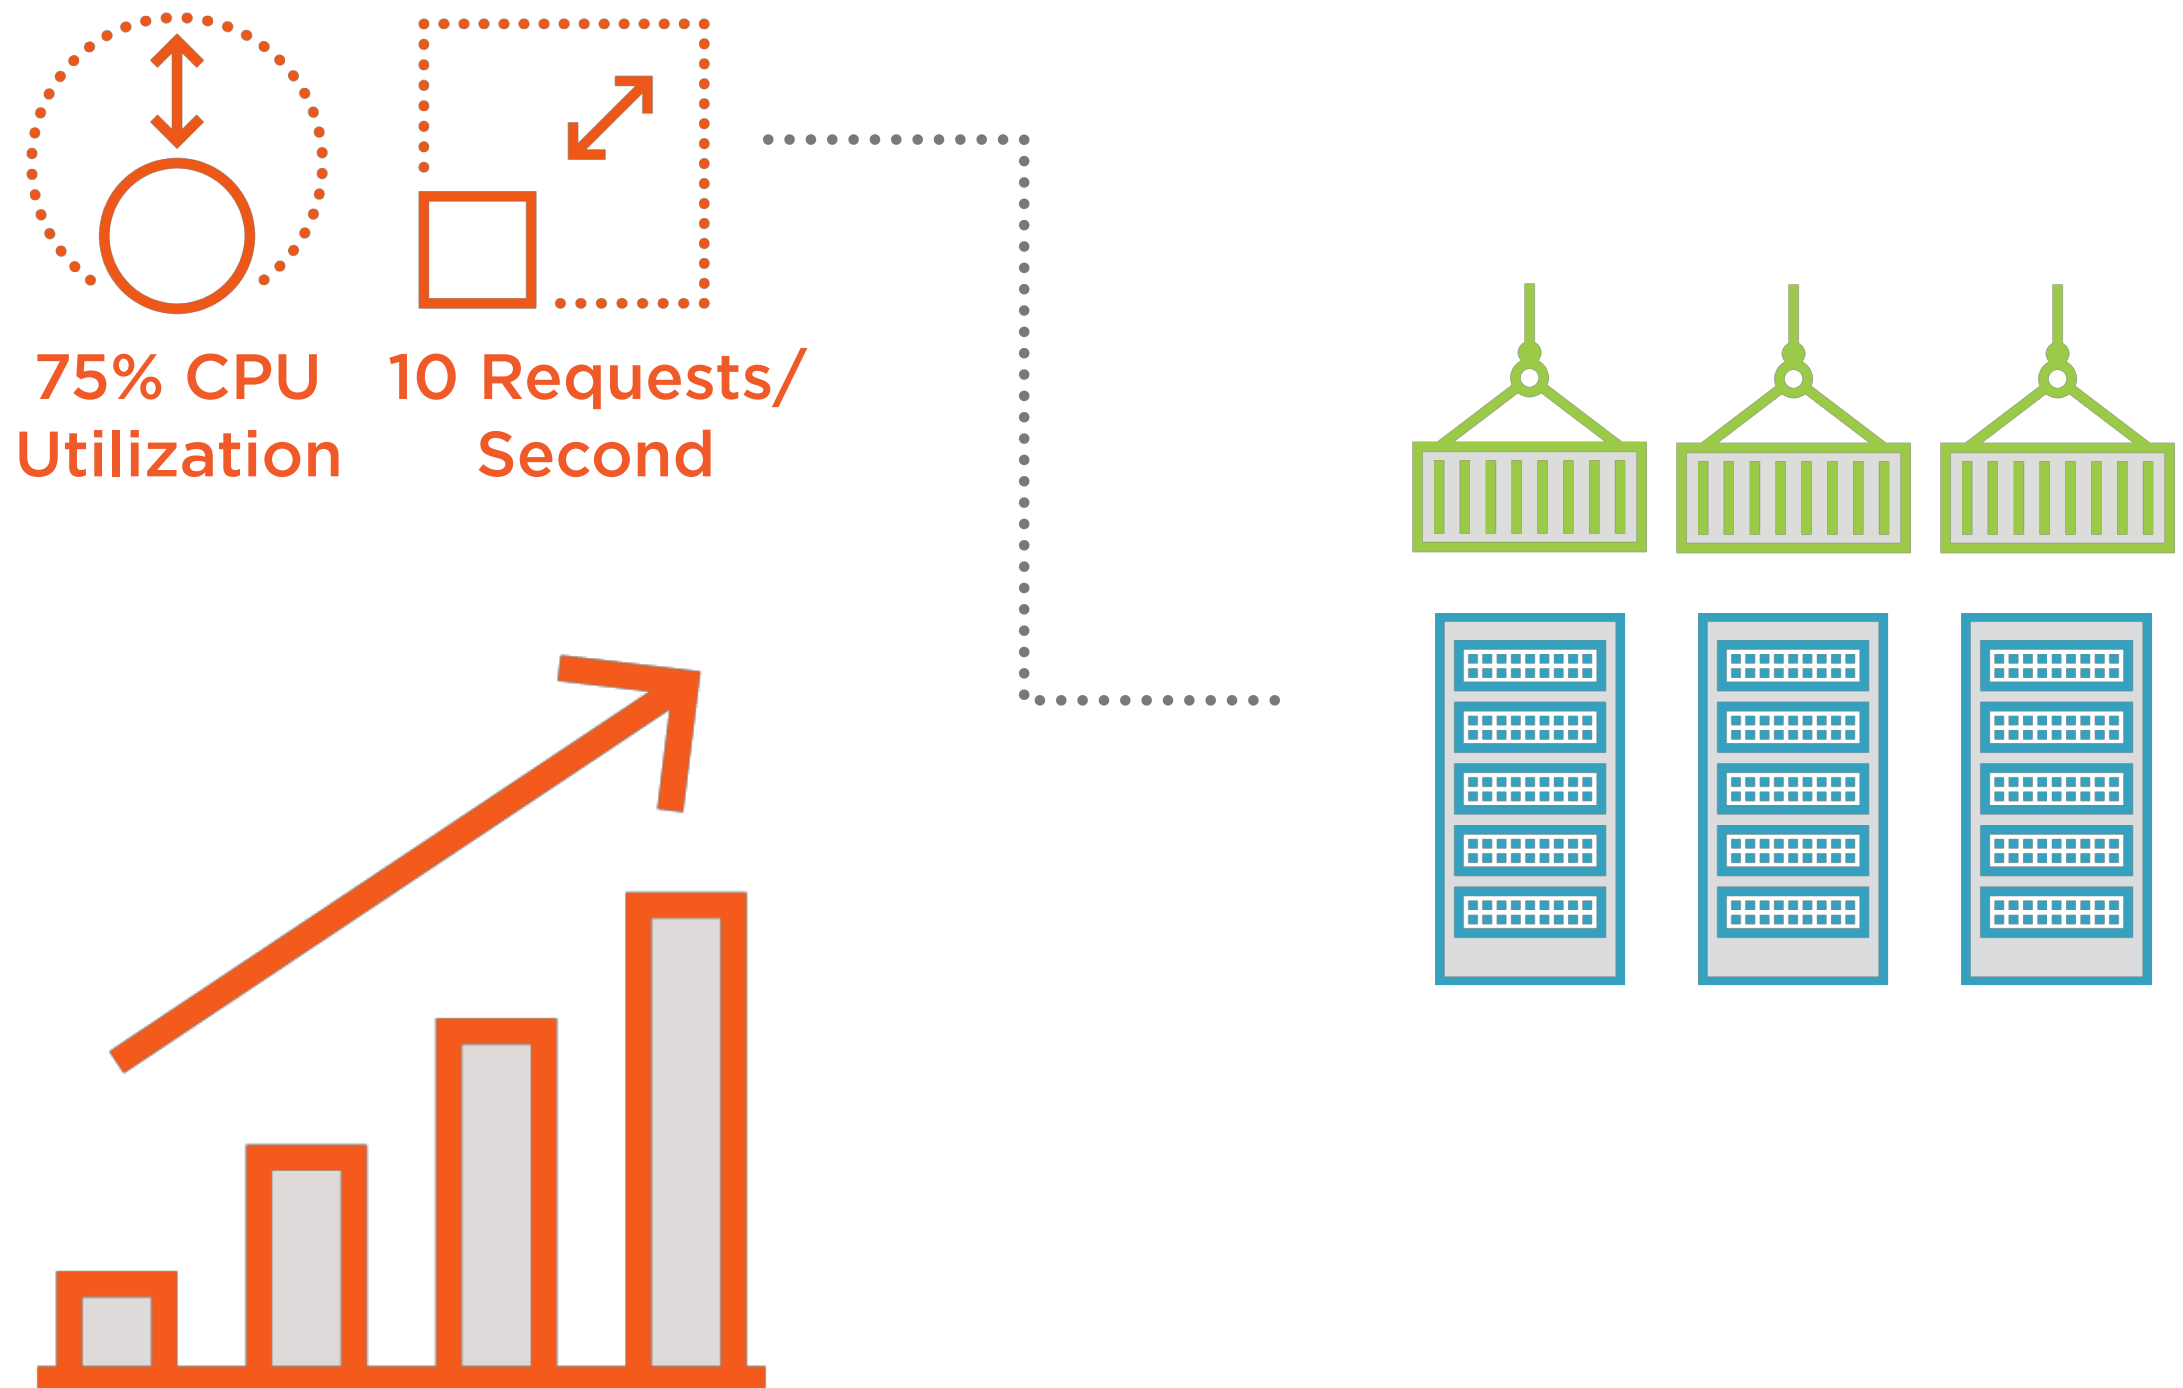
Platform ensures health of compute resources



# Autoscaling with Kubernetes




# Autoscaling with Kubernetes






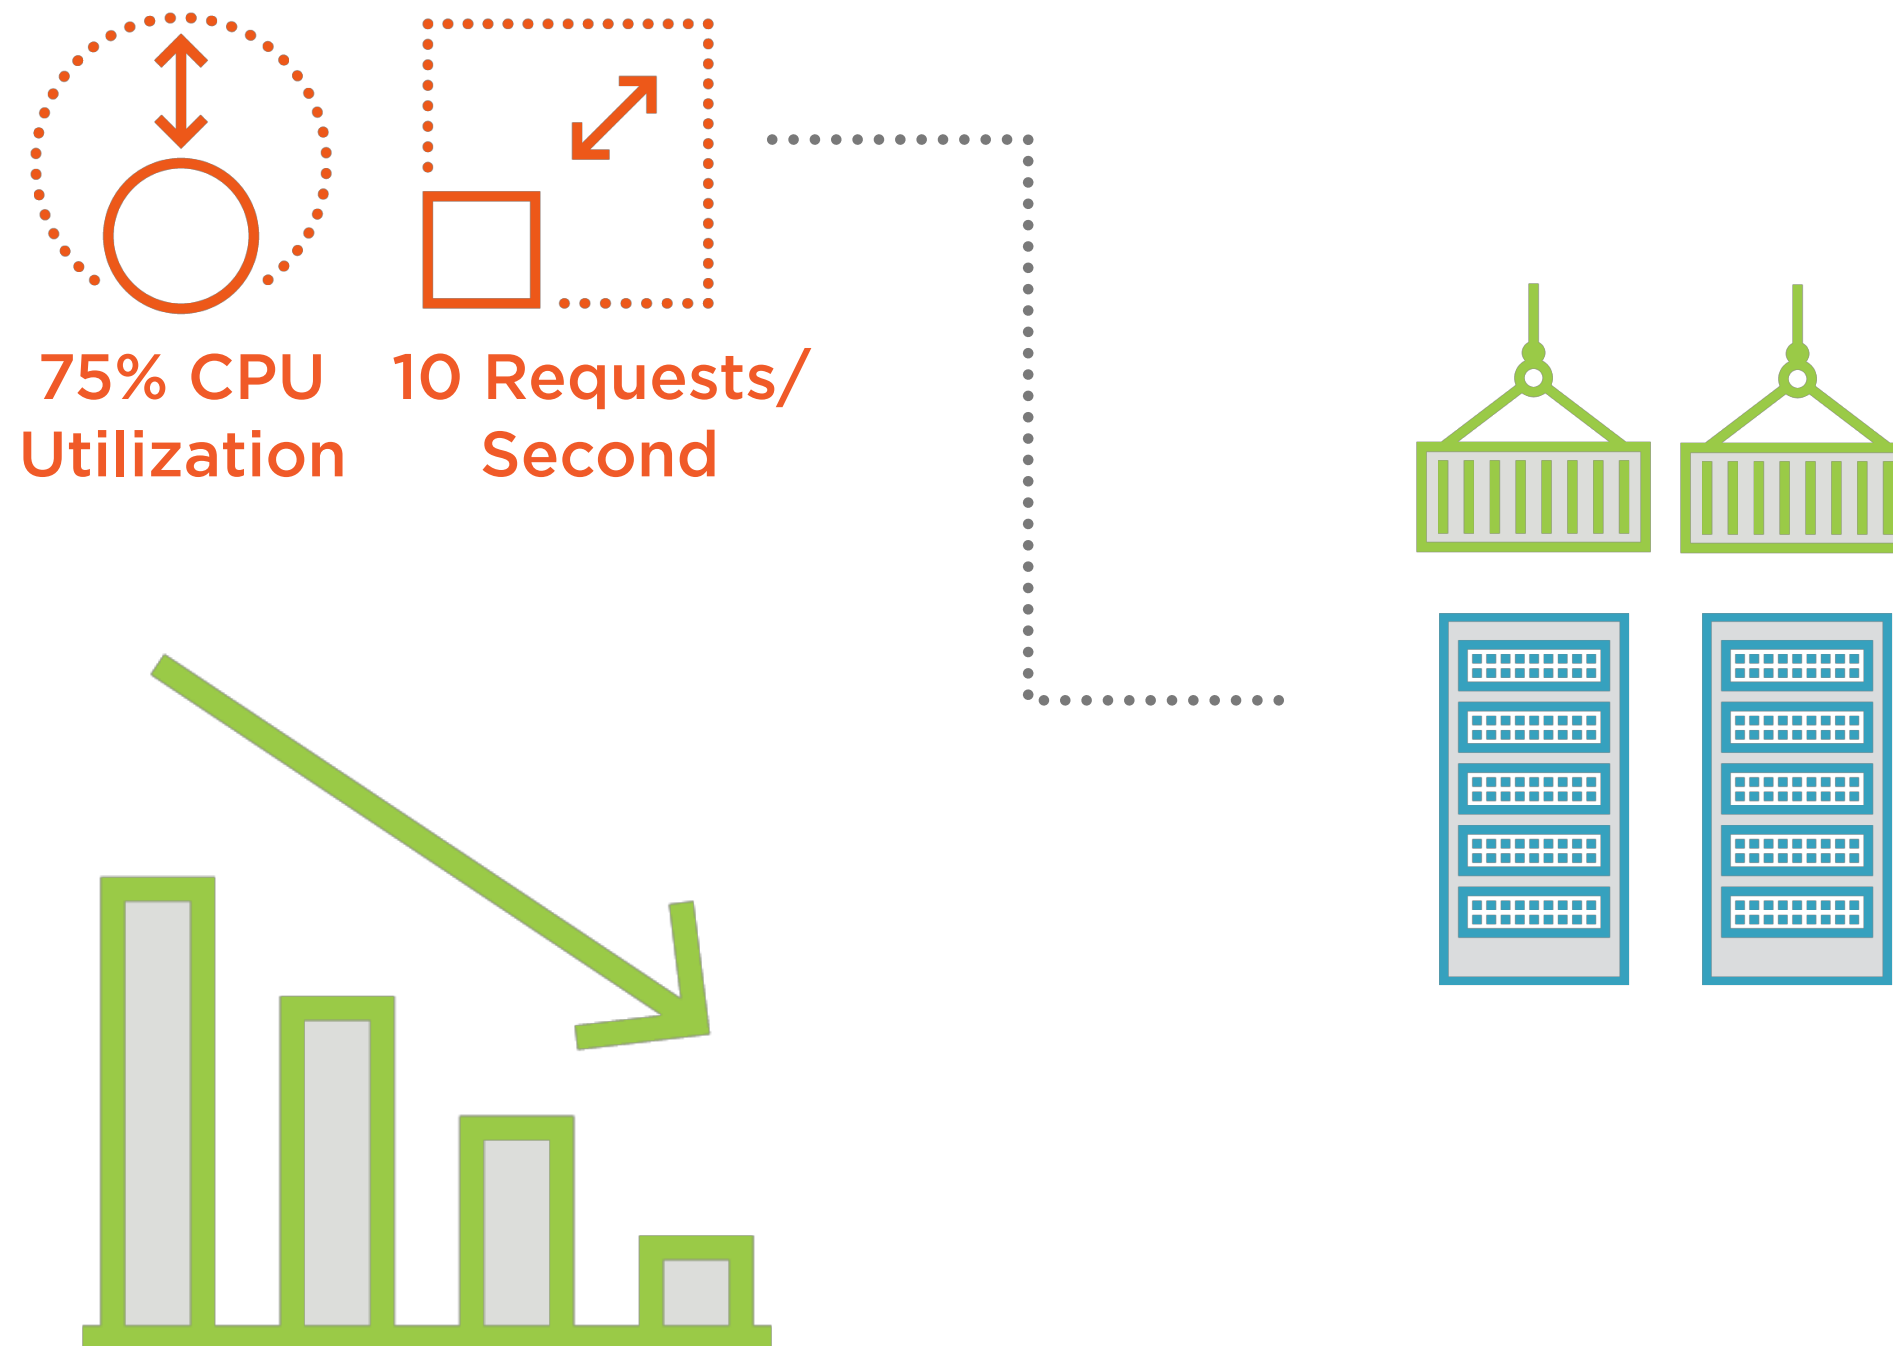
# Autoscaling with Kubernetes



# Autoscaling with Kubernetes



# Autoscaling with Kubernetes





# Horizontal Pod Autoscaler

**Even higher level abstraction**

**Specify any scalable object as target**

**Along with autoscaling policy**



# HPA Targets



## Scalable objects

- ReplicaSet
- Deployment

## Can't target non-scalable objects

- DaemonSets



# HPA Metric

**CPU utilization**

**Custom metrics**





# Preventing Thrashing

**Thrashing occurs when HPA can't find stable size**

**Use cool-down periods to avoid this**


**Set intervals between successive operations of same kind**

# GKE Pricing

---







# Compute and Storage

**GCE instances are used for nodes in the cluster**

**Billed for instances based on GCE usage**

**Per-second basis with 1 minute minimum cost**

**Costs vary based on machine type, region, type of disk etc.**



# Summary

**Containers for lightweight compute**

**Ideal for hybrid, multi-cloud**

**Kubernetes is a container orchestration technology**

**Industry standard with Google origins**

**GKE for Kubernetes on GCP**

